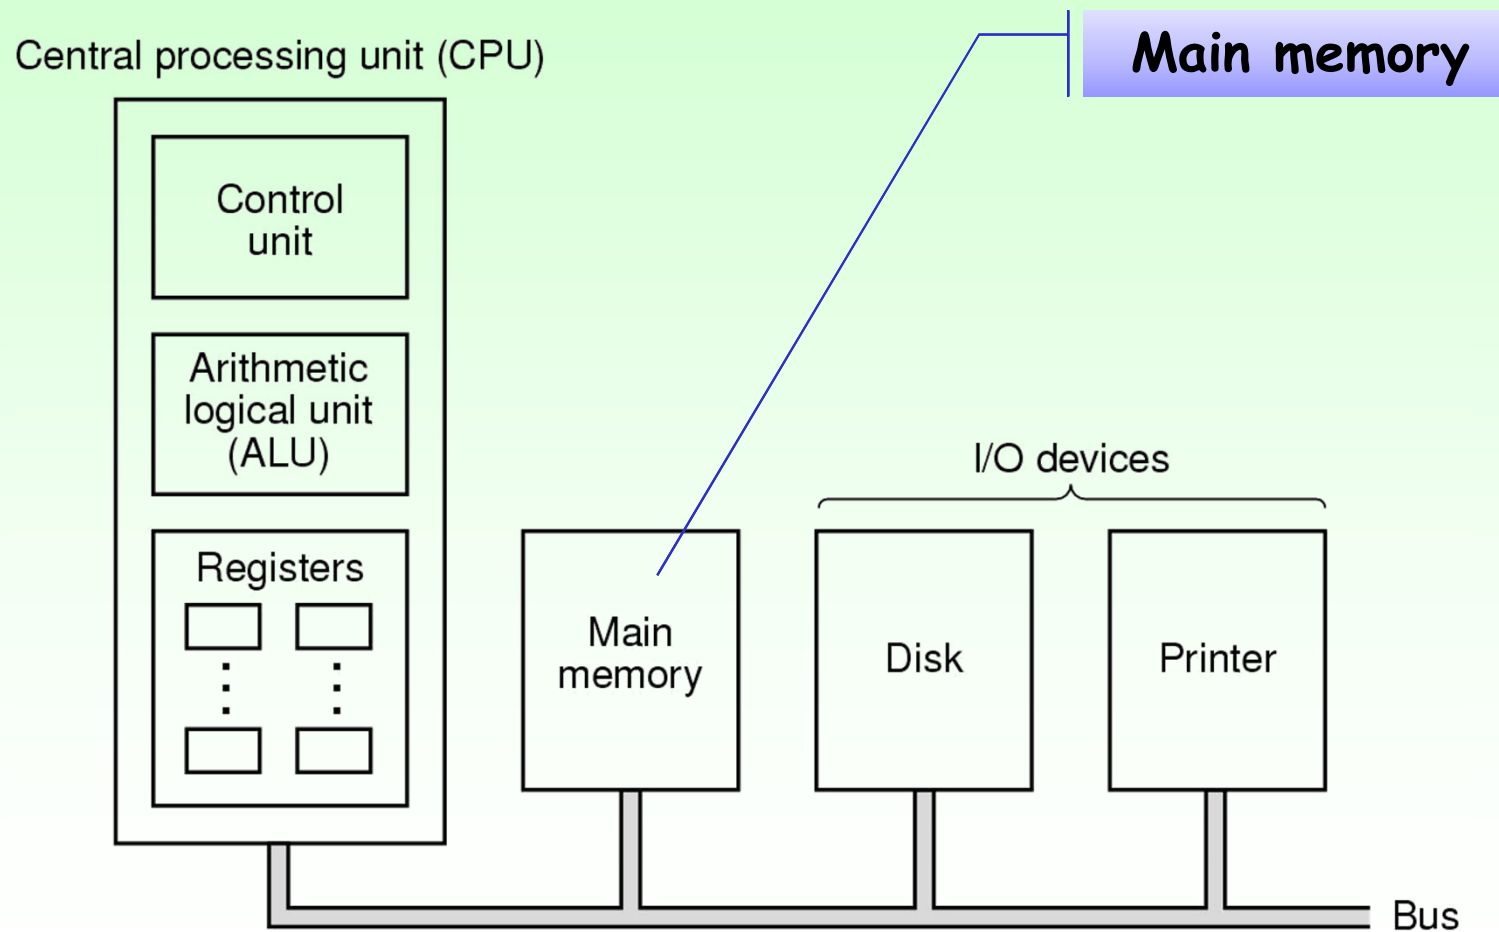


# CSCU9V4 Systems

Systems Lecture 7  
Computer Organisation 2

**Main Memory & Cache**

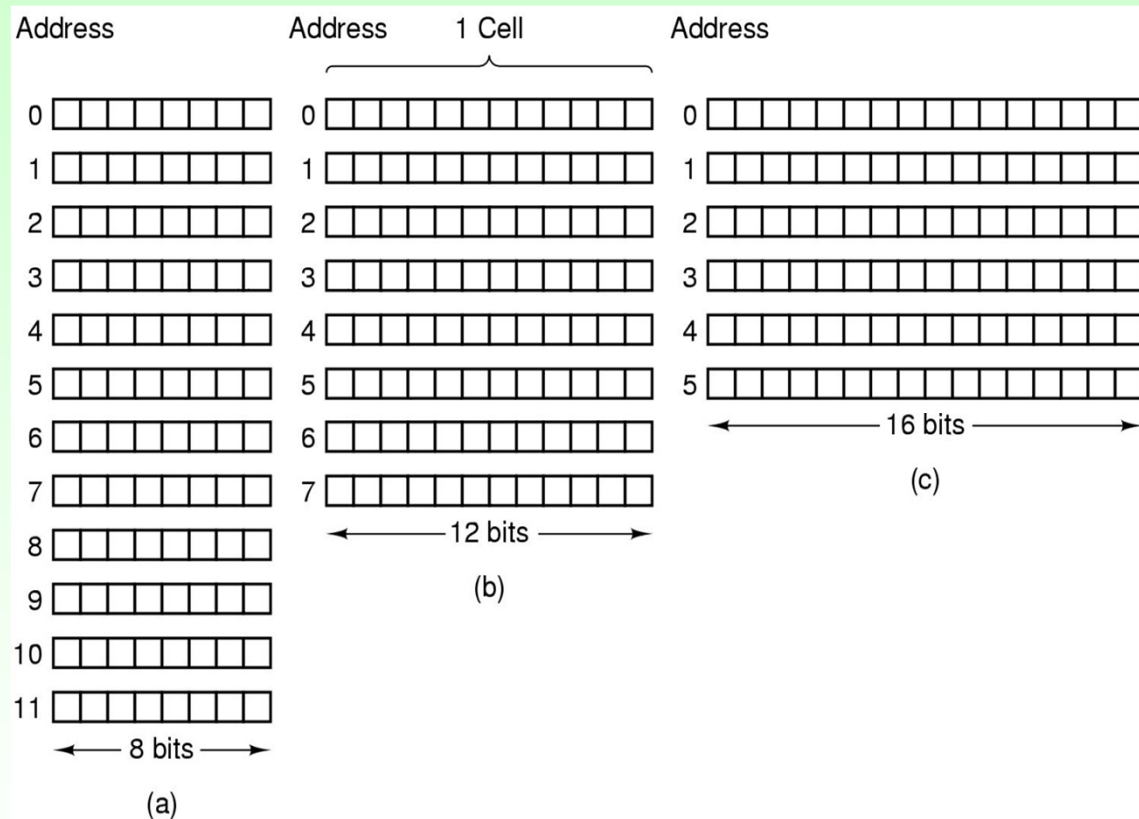
# Structure of a Simple Computer



# Memory Cells

To the programmer, the main memory appears as a 1-dimensional array of *cells*.

- A *cell* is the smallest element of memory that can be accessed.
- Every cell has a unique address, starting (generally) at 0.
- Each cell is individually addressable.
- If a memory has  $n$  cells, their addresses will usually be  $0 - (n-1)$
- All cells in a computer are the same size (hold the same number of bits)
- Example shows 3 ways of organising cells
- Adjacent cells have consecutive addresses



# Addressing Memory

- Computers use the binary system therefore express memory addresses as binary numbers
- An address with  $m$  bits, can be used to address  $2^m$  memory cells.
- Previous example:
  - a) requires 4 bits ( $2^4 = 16$ );
  - b) requires 3 bits ( $2^3 = 8$ );
  - c) requires also 3 bits, because  $2^2 = 4$  which is too small
- The number of bits in an address determines the number of directly addressable cells; it is independent on the number of bits in a cell
  - $2^{12}$  cells with 8 bits each and  $2^{12}$  cells with 16 bit each, both need 12-bit addresses
- Cell is the smallest addressable unit!
  - Generally a byte: also explains why memory is sold by the byte (or Mbyte, Gbyte, or Tbyte)
  - Also explains why memory capacity is a power of 2.

## Example Word and Cell Sizes


<u>Processor (microprocessor)</u>	<u>Cell Size</u>	<u>Word Size</u>
6800, 6502, 6809, 6802, 8088	8	8
68000, 68010, 68020, 68030, 68040	8	32
8086, 80286	8	16
80386, 80486, Pentium-Pentium 4	8	32
PowerPC G5, DEC Alpha, recent Intel processors (IA64, Core Ix, Nehalem, Sandy Bridge, etc)	8	64

### Addressable memory:


- Size of MAR defines number of addressable cells.
- 16 bits  $\Rightarrow 2^{16}$  cells = 64K cells (usually 64K bytes) (early microprocessors)
- 32 bits  $\Rightarrow 2^{32}$  cells = 4G cells (usually 4G bytes) (Pentium series)
- 42 bits  $\Rightarrow 2^{42}$  cells = 4T cells (usually 4T bytes) (Power PC G5, Intel Xeon E7 V2)
- Word size and address size now often the same eg 32 or 64-bit

## Aligning Words

- Many architectures require words to be aligned to their natural boundaries (8-byte word may begin at address 0, 8, 16, ...; but not at 4, 12, or 19)
- Memories operate more efficiently if they are aligned
- Pentium fetches 8 byte at a time, uses 36-bit physical address space, but only has 33 address bits
  - 3 least significant bits are not specified
  - No non-aligned memory reference possible!



(a)



(b)

## Problems with Aligning Words

- Pentium II needed to be able to access memory at any location
  - Two memory accesses needed
  - Example: 4-byte word at address 7 (bytes 7 - 11)
  - One reference to address 0 (get bytes 0 - 7)
  - A further reference to address 8 (get bytes 8 - 12)
  - Hardware extracts the relevant bytes and puts a 4-byte word together
- Required extra logic on the chip -> slow and expensive!

## Intel's advice

- **Intel® Pentium® 4  
and Intel® Xeon™ Processor Optimization  
Reference Manual**
- For best performance, align data as follows:
- Align 8-bit data at any address.
- Align 16-bit data to be contained within an aligned four byte word.
- Align 32-bit data so that its base address is a multiple of four.
- Align 64-bit data so that its base address is a multiple of eight.
- Align 80-bit data so that its base address is a multiple of sixteen.
- Align 128-bit data so that its base address is a multiple of sixteen.



# Memory Performance


- CPUs are much faster than memory
  - techniques such as pipelining, superscalar architectures make the CPU's memory demand very high indeed
  - A 2 GHz CPU (with a pipelined superscalar architecture) may need up to 8,000,000,000 instructions every second
  - A 2 or 3GHz processor is limited by operating over a 1GHz data bus
- Problem: Supplying sufficient data and instructions to the CPU
  - The “Von Neumann Bottleneck”

# Improving Performance

- Faster memory?
- Use even faster buses (connection between main memory and CPU)?
- Place ultra-fast memory on the CPU chip
  - On-chip memory can be accessed much faster (why? nearer, + arguments from analogue signal processing, not appropriate for this course)
- Use cache memory
  - High speed memory on (or very close to) CPU chip
- Combine fast and slow memory
  - Aim to get the speed of the fast memory and the size of the slow memory
  - Keep the most heavily used words in the cache memory
  - If the CPU needs a word it first checks the cache
  - Which data should be stored in the cache?

# Cache Memory

- Cache memory
  - Usually placed on the CPU chip
  - access is very fast
    - 10-100 times faster than main memory
  - memory bus is used for
    - cache misses
    - when data is written (need to maintain consistency).



# Cache Misses

- A cache miss occurs when the CPU needs an item (instruction or data) which is not currently in the cache.
- When this occurs action must be taken:
  - item must be copied from main memory into the cache,
  - replacing something which is already in the cache
    - which may have to be saved in main memory before it is overwritten in the cache
- There are complications here:
  - which item in the cache will be selected to be overwritten?
- But it's worth it ...

## Why does Cache Memory Help?

- Cache memory helps in practice because programs exhibit *locality of reference*.
- That is, recently used instructions and data are likely to be used again
  - true both for programs and data
  - particularly helpful for programs, because programs are read-only
    - no problems with writing and consistency.
- In particular, programs contain loops
  - most programs spend a lot of their time executing loops, going over the same instructions that are likely to be in the same area of memory
  - the larger the cache, the more likely the required data/instructions will be within the area copied to the cache

## Example

```
for (index = 1; index<MAXXPOINTS; index++) {  
    endpoint = floor(index *  
        ((float)displayframelength/MAXXPOINTS));  
    /* find mean of points */  
    pval = 0;  
    for (j = startpoint; j<endpoint; j++)  
        pval += displayframe[j];  
    pval = pval/(endpoint - startpoint);  
    startpoint = endpoint;  
    Xlocn += Xinc;  
    PSlineto(Xlocn, pval * Ymag + Yaxis);  
}
```

Say MAXXPOINTS is 250, and endpoint might be 50 to 500 each time: how much time is spent in the innermost loop? )

## Example (Continued)

- This fragment of code (from a real program) is generally executed thousands of times. The code will all fit into a cache of any reasonable size.
- It will therefore only be fetched once - so relieving the CPU/Memory connection of a considerable amount of work.
- Many data items change each iteration - but not all of them. Here

`displayframelength, Ymag, Yaxis, Xinc`

are all the same in each iteration.

- The use of a cache memory can make a huge difference to the amount of data that needs to be transferred from the main memory to the CPU.
- The overhead of the cache is to be found in
  - greater circuit complexity,
  - dealing with cache misses.
- Cache hit rates are usually 90-99% for a reasonably large cache

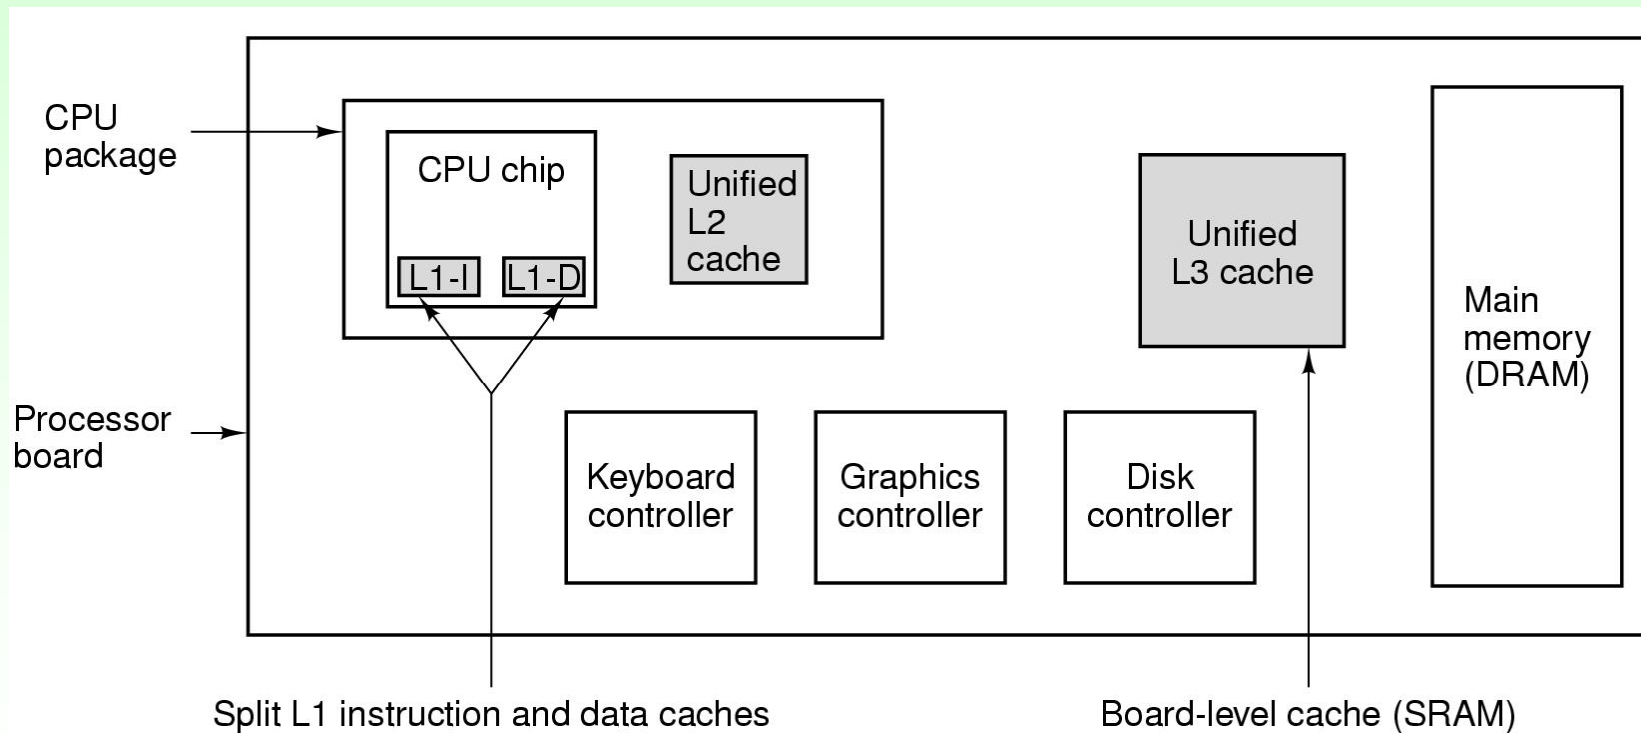
# Design of Caches

- Locality principle is used as a guide
- Main memories and caches are divided into cache *lines*
  - E.g. 64 byte line
- A full cache line is loaded into the cache
- More efficient than loading individual words
- Cache design is important:
  - Cache size
  - Size of cache lines
  - Cache organisation
  - Data and instructions kept in the same cache?
  - Number of caches?



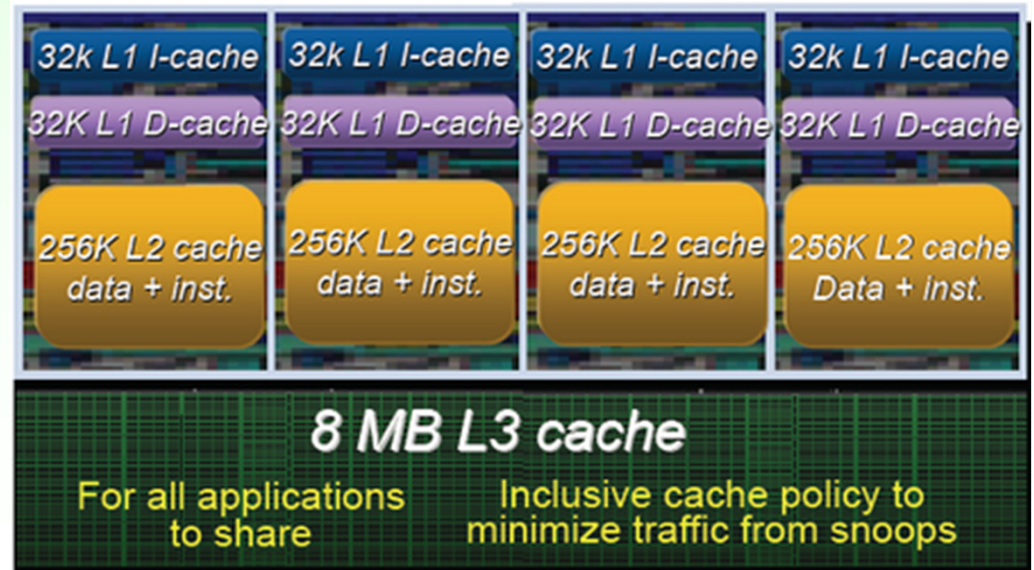

# Number of Caches

- Latency (time to load a memory item)
  - can be improved by a reasonable sized cache
- Bandwidth can be improved by having multiple caches



## Multi-level caches

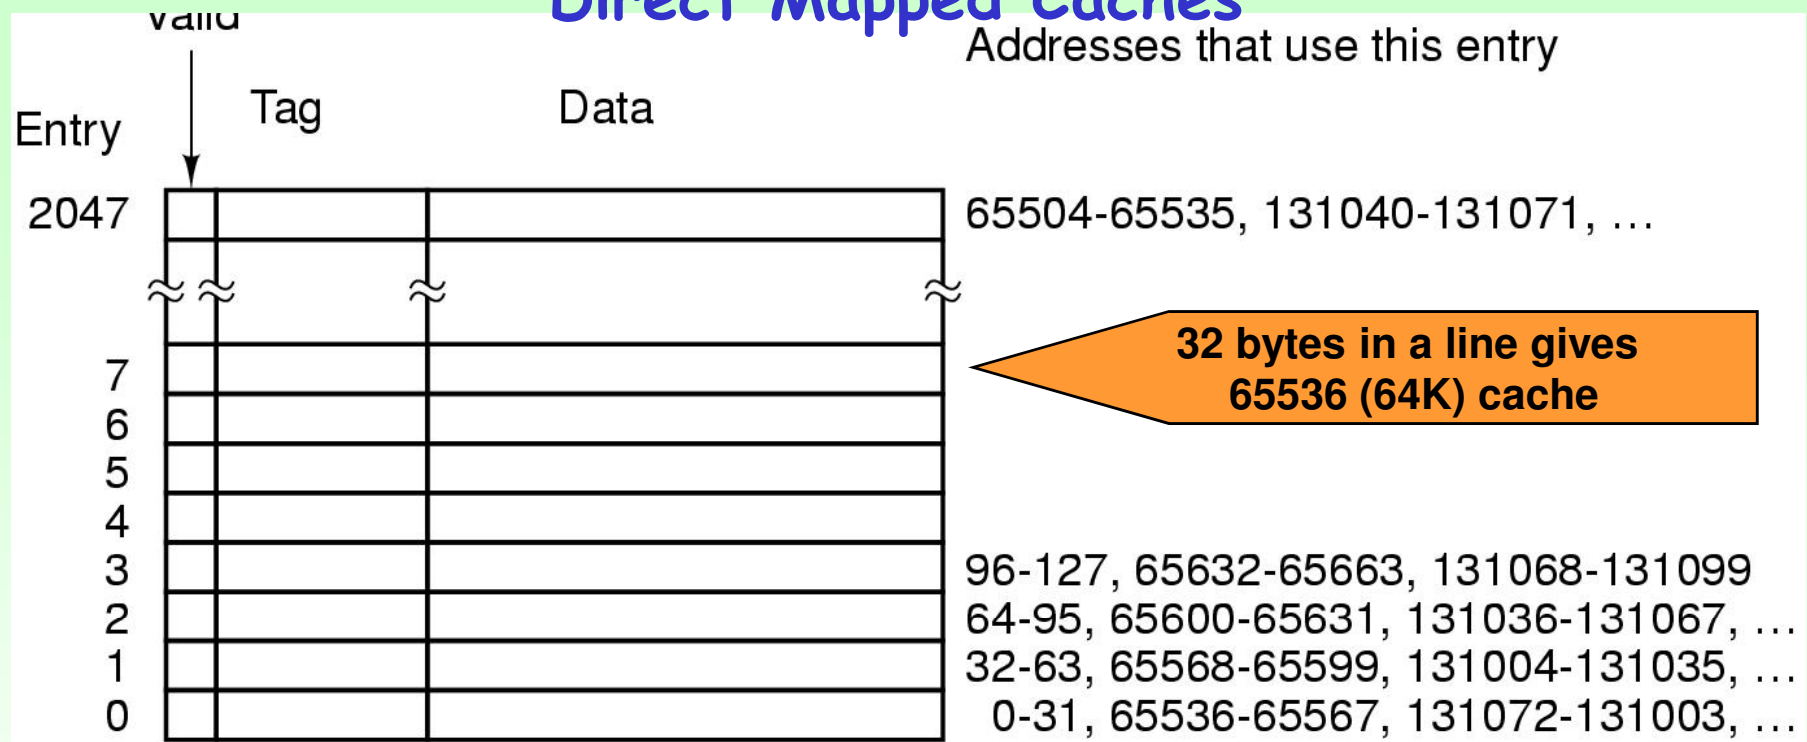
- Modern CPUs use multi-level caches, on-chip.
- Example is from Intel i7.
- L1 cache is in 2 parts, an instruction cache and a data cache: one per core
- L2 cache is unified: one per core
- L3 cache is unified across all (4) cores.



# What to keep in the cache?


- Spatial locality
  - Words located close to the currently required one
    - Expected to be used soon
- Temporal locality
  - Words most recently used
    - Likely to be used again soon
- There are many forms of cache memory:  
Next few slides show some of these: included for completeness, not examinable!

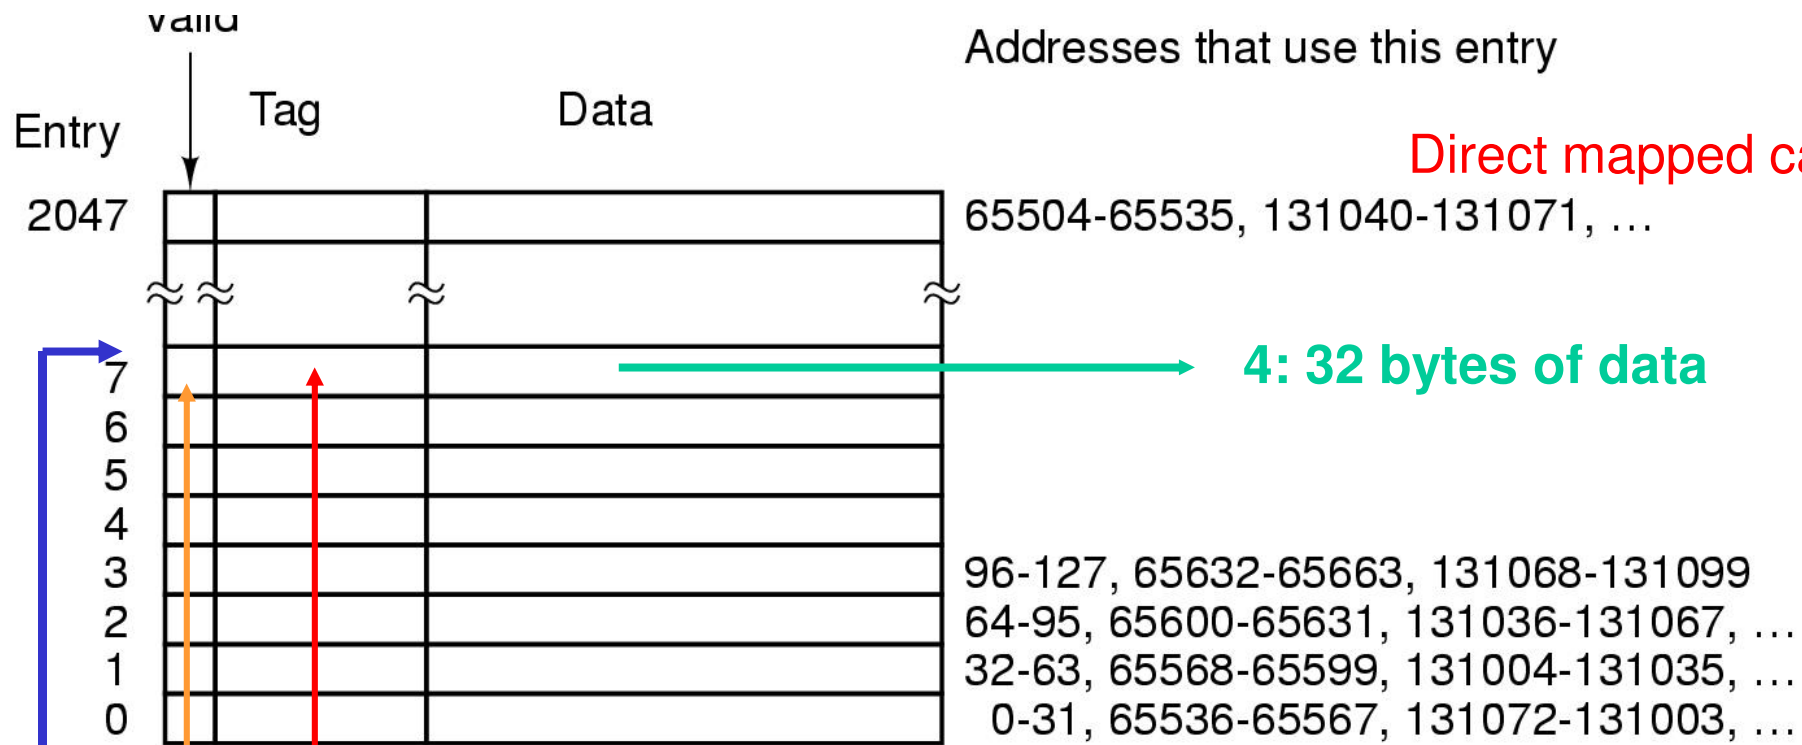
# Direct Mapped Caches




Direct mapped cache (a)

A 32 bit address

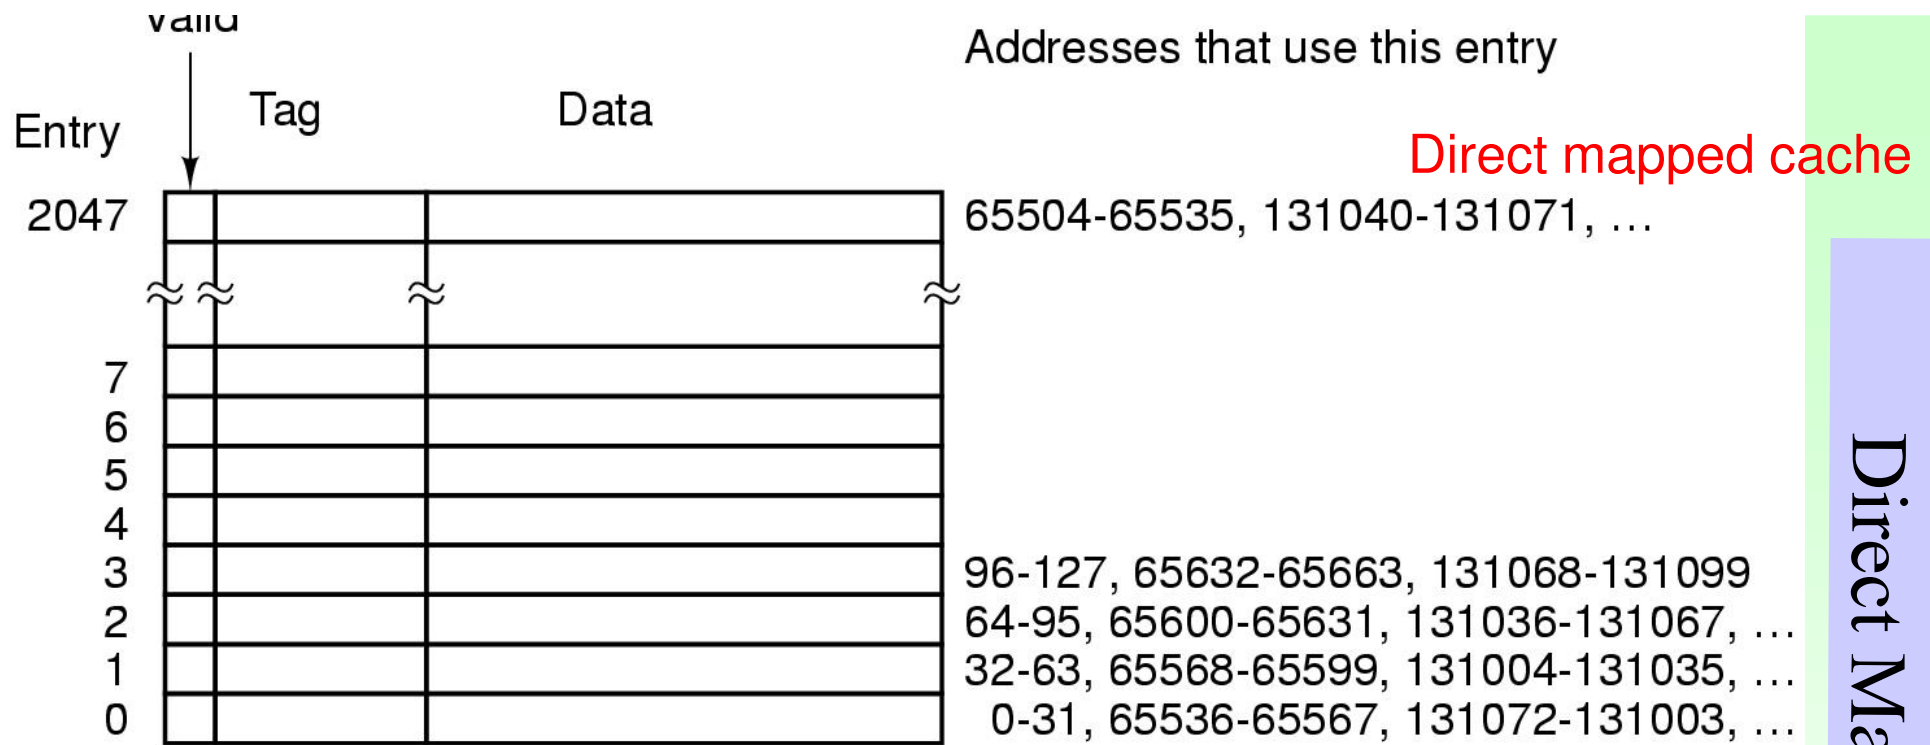





(a)



# Direct Mapped Caches



(a)

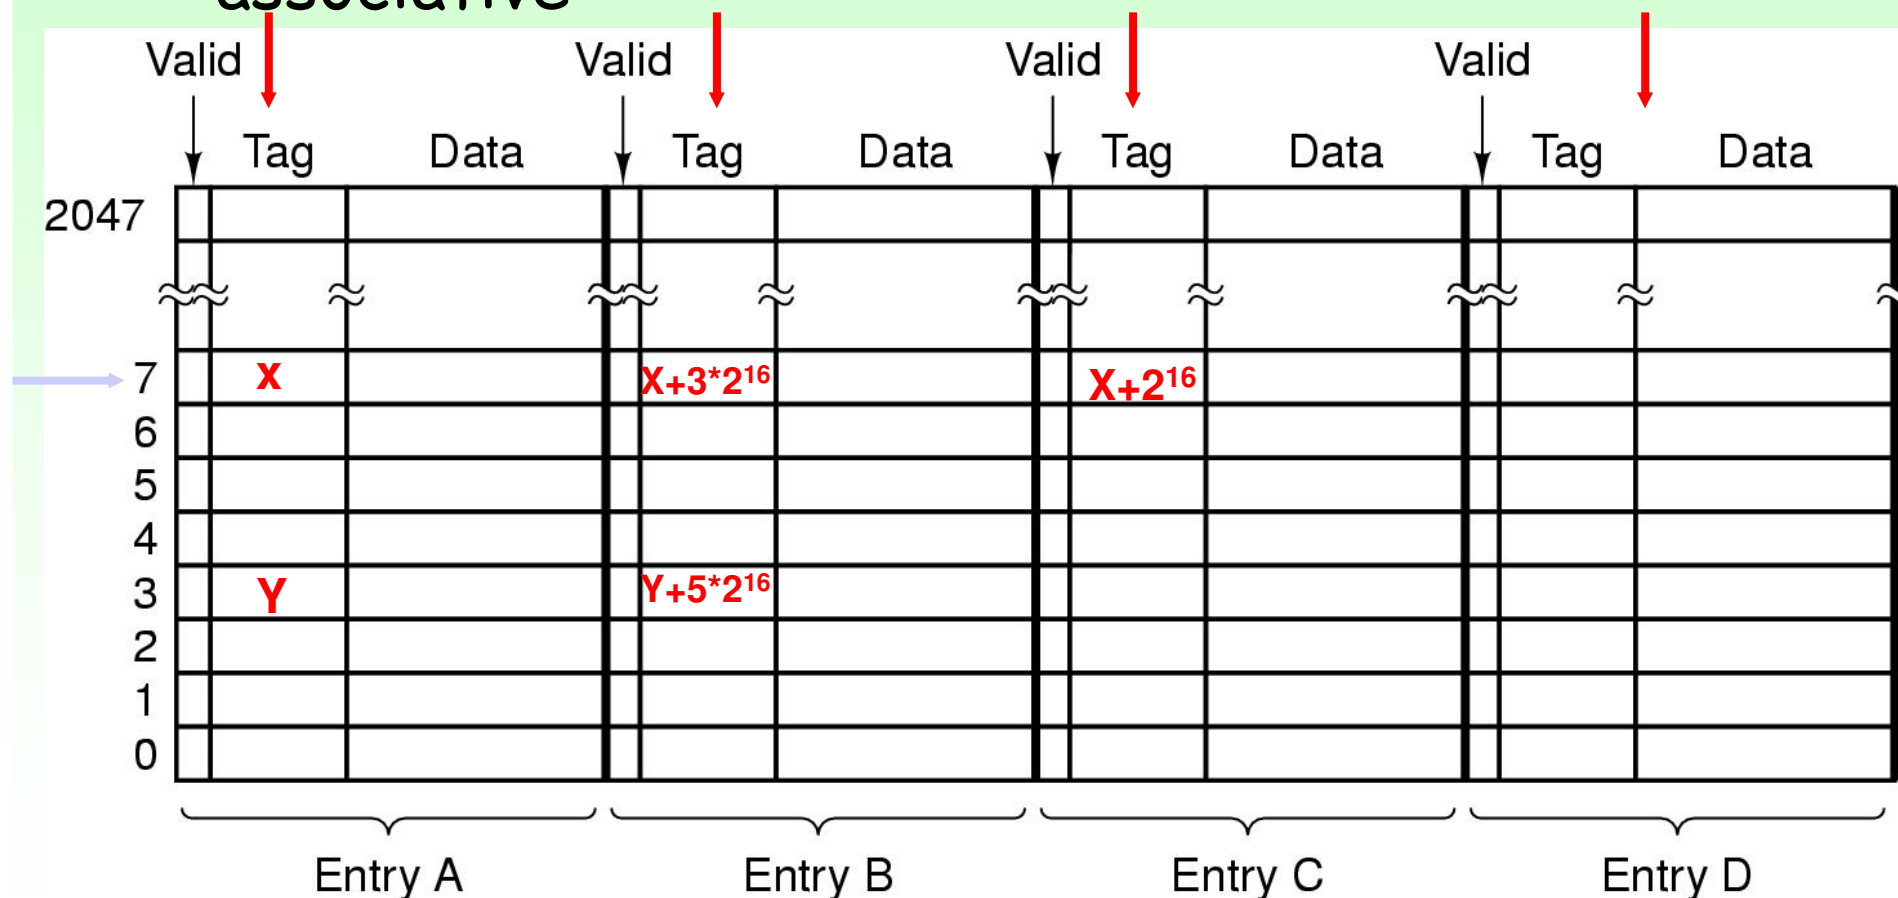


Direct Mapped Caches

## Set Associative Caches

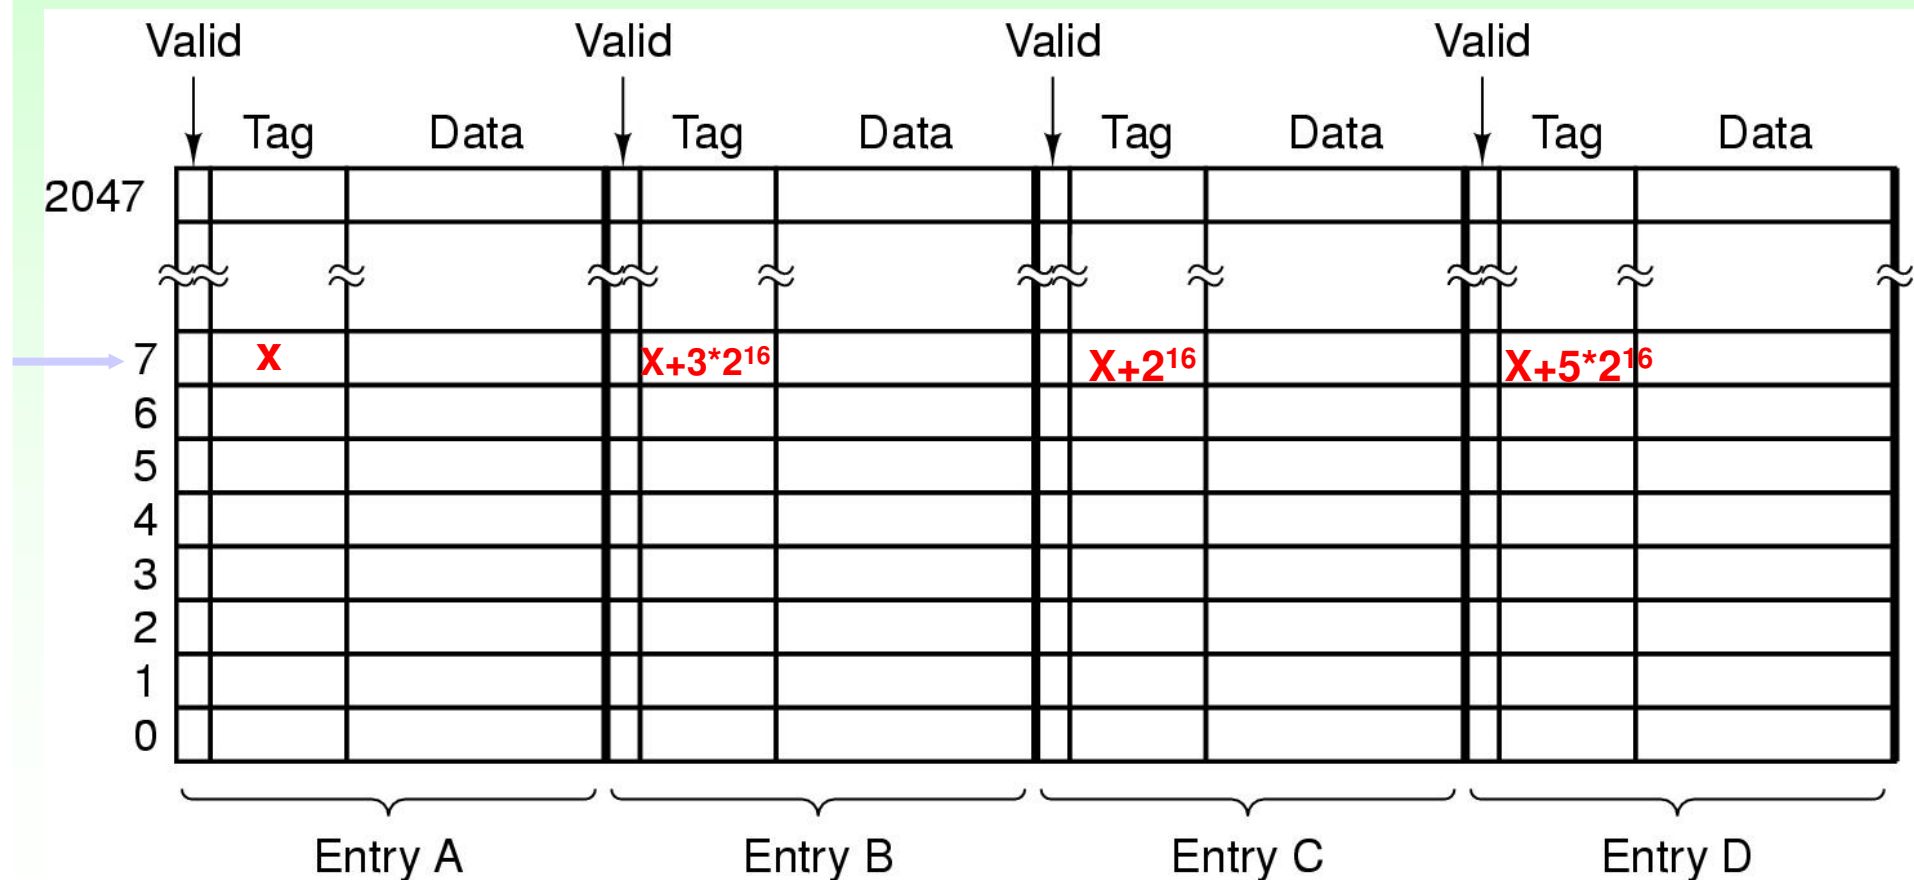
- Allows two or more cache lines in each
- A cache with  $n$  lines per entry is called  $n$ -way associative

$n$  is  
normally  
2 or 4



## Set Associative Caches

- What if a line is full (all  $n$ -entries used) ?
- Overwrite Least Recently Used - LRU





# Write operations : a special problem for caches

- Write a word
  - Either update the cache or make entry invalid
  - Usually entry is updated
  - What about main memory?
- Write immediately = **write through**
  - Simpler to implement
  - More reliable -> memory is always up-to-date
  - Consider errors -> recover the state of memory

Entry	Tag	Data
2047		
~	~	~
7		
6		
5		
4		
3		
2		
1		
0		

- Write later (on re-use of Cache line) = **write deferred**
  - Used by more sophisticated implementations
- Writing data which is not in cache - do we load it into cache?
  - Yes : most "write deferred" implementations use **write allocations**
  - No : "write through" implementations would complicate a simple design

**write miss**

**End of Lecture**