

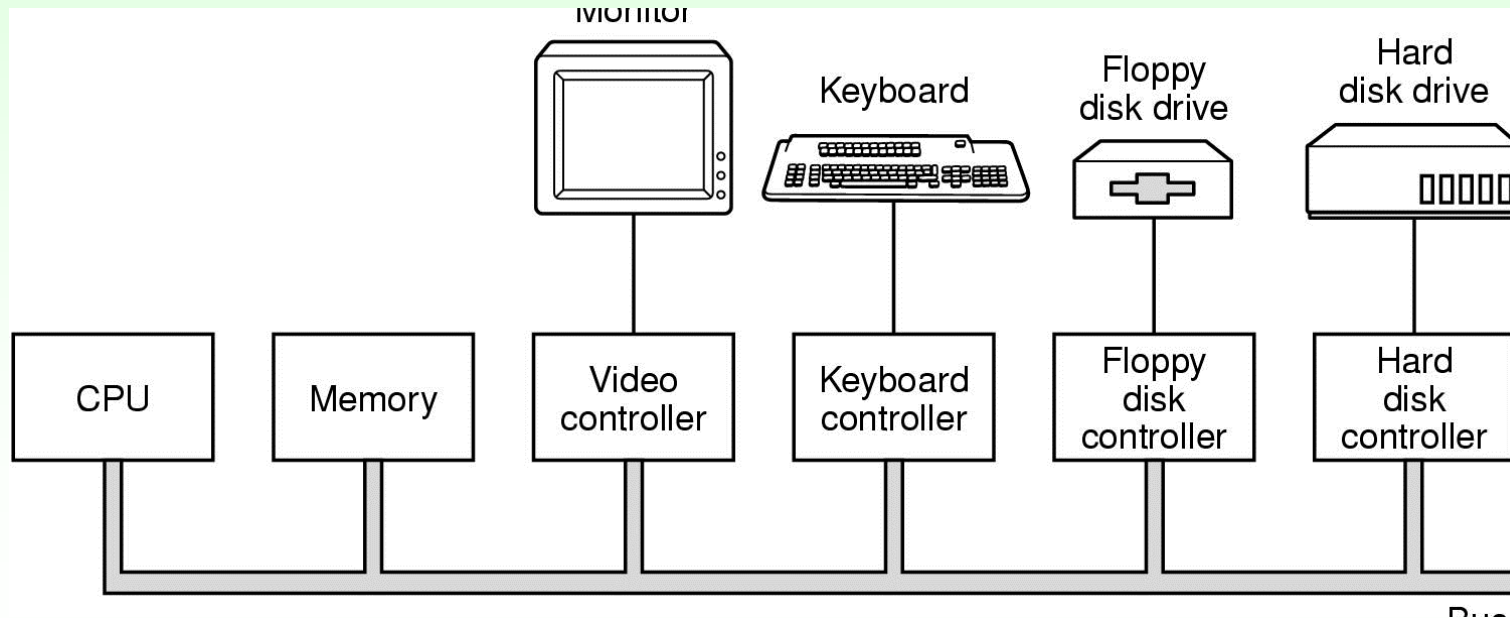
CSCU9V4 Systems

Systems Lecture 8
Computer Organisation 3

Buses

Buses

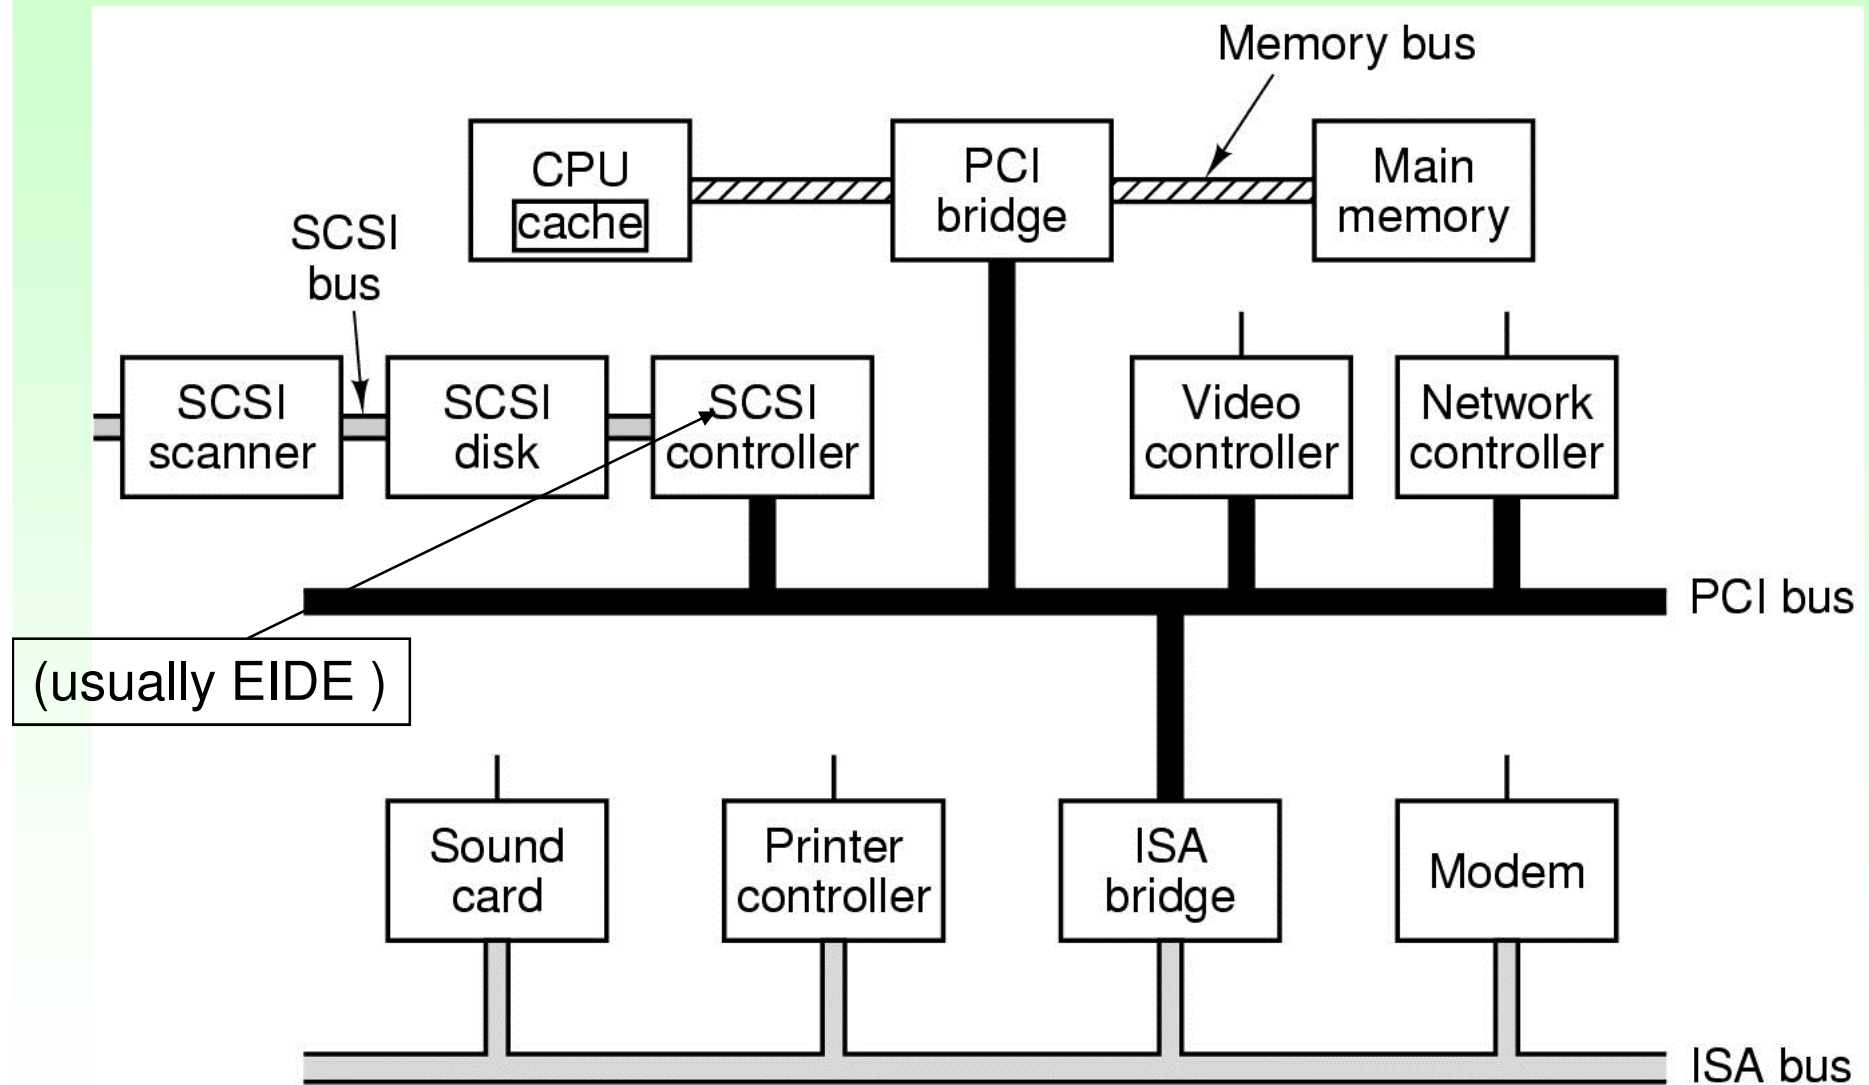
- Origin: same as omnibus - meaning 'for all' (Latin)
- High-speed data link between all the devices attached to the bus.
- Connect I/O controllers to the memory
- Used by the CPU to access memory
- Possible conflict if two (or more) devices and/or the CPU want to access the bus - requires Bus Arbitration to choose Bus Master



Development of Buses

- Single bus design was common for early computers and worked fine
- Usually a faster bus was designed for the next system
- Developing a new bus for every new PC model is impractical
- Old IBM PC bus is called ISA (Industry Standard Architecture)
- Machines with multiple buses were developed
 - CPU is often connected to memory by a dedicated high-speed bus
 - So that CPU traffic does not affect other buses
 - PCI (Peripheral Component Interconnect) used for most other devices
 - High-bandwidth peripherals are connected to the PCI bus
 - ISA or EISA bus can be used by older devices (may be absent entirely in new machines)
 - PCI has a connection to the ISA bus if it is present
 - Modern machines use SATA and PCIe.

Example Architecture (old and simple)



Different Buses

- Buses are inside a shielded case to minimise electrical interference
 - Transmission is error-free
- Buses inside the CPU may be of any kind (and are not studied here!)
 - Do not interface with the outside world
- External buses must adhere to well-defined rules on how the bus works
 - Allow 3rd parties to develop peripherals and controllers
 - Rules are called the bus protocol
 - Additionally there are electrical and mechanical specifications
- Some example buses:
 - Multibus (8086)
 - IBM PC bus (PC/XT)
 - ISA bus (PC/AT)
 - EISA bus (80386)
 - Microchannel (PS/2)
 - SCSI (PCs and workstations)
 - Nubus (Macintosh)
 - Universal Serial Bus (PCs)
 - FireWire (consumer electronics)
 - VME bus (physics lab equipment)
 - PCI bus (PCs)
 - PCI Express bus (newer PCs)


Master and Slave

- Some devices attached to buses are active and can initiate transfers
- Other devices are passive and wait for requests
- Active devices = Masters
- Passive devices = Slaves
- Example:
 - CPU is bus master, and orders the disc controller to read a block
 - Disc controller becomes bus master, and transfers the block from disc to memory
- Master and slave associations may change and depend on particular activities, however, memory can never be master

Master	Slave	Example
CPU	Memory	Fetching instructions and data
CPU	I/O device	Initiating data transfer
CPU	Coprocessor	CPU handing instruction to Coprocessor
I/O device	Memory	Transfer data to memory
Coprocessor	CPU	Coprocessor fetching operands from CPU

Bus Outline

- The bus consists of a number of parallel conductors
- These are used to specify
 - Memory address (or peripheral being accessed)
 - Direction the transfer is running in
 - Timing of the transfer
 - Data to be transferred



From Wikipedia

- Often separate sets of conductors are used for these various purposes:

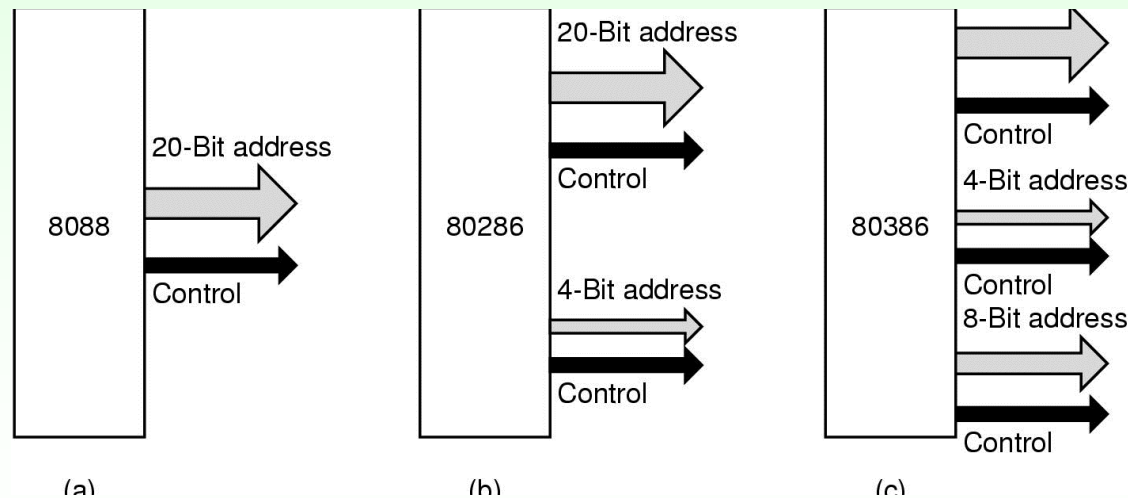
Line Name	Symbol	Function
Address lines	$A_0..A_{NA}$	Specifying the cell being accessed (same width as MAR)
Data lines	$D_0..D_{ND}$	The data being transferred (same width as MDR)
Read/Write	R/W	Direction of transfer (wrt CPU)
Clock	f	Timing of the transfer

Bus Outline

- Bus design is a complex problem
- We will provide an overview here
- Principal bus design issues include
 - Bus width
 - Bus clocking
 - Bus arbitration
 - Bus operations
- Look at them in turn ...

Bus Width - Address lines

- The more address lines are available the more memory can be addressed directly
- Bus with n address lines \rightarrow CPU can address 2^n different memory cells
- Large memories \rightarrow many address lines
 - 32 bits: 4Gcells (usually 4Gbytes)
- Wide buses require more wires and also take up more space
- More expensive! \rightarrow memory size vs. cost
- Needs to be considered if the memory is to be extendible later



Bus Width - Data lines

- How to increase the data bandwidth of a bus
 - Use more data lines
 - Decrease the bus cycle time (more transfers/sec)
 - Increase the data bus width (more bits/transfer)
- Speeding up the bus is difficult (signals travel at different speeds on different lines)
 - Electrical problems increase with increasing bus speed
 - Old devices can't work at higher speeds
- Hence, usual approach to improve performance is to increase data width
 - Add more data lines
- Another option: multiplexed busses (use signal lines for more than one purpose)
 - Cost vs speed



Mode of Transfer (Control Lines)

- Read or Write
 - R/W line
- There may be additional lines specifying the size of the transfer:
 - the bus can transfer 1 word (i.e. using all data lines) at a time, but sometimes it may transfer less:

8 bits	Byte
16 bits	Short
32 bits	Word
64 bits	Long word
- Some buses also have 'burst mode' transfers:
 - for example, several contiguous words of memory may be transferred, one after the other from memory to CPU
- Interrupts...see later

Timing of Transfer

The control signals specify the times ...

... when the processor (or memory or peripheral devices)

- may place data on the bus,

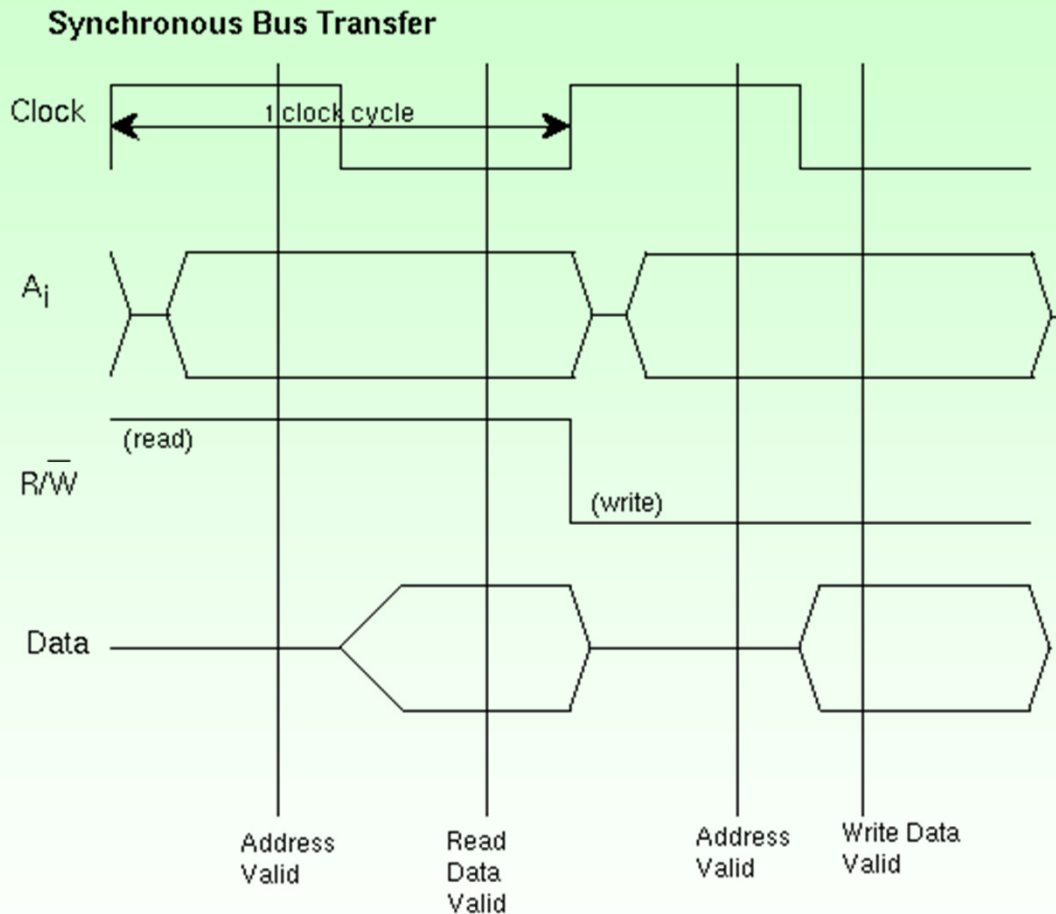
... or

- when the data on the bus may be assumed to be valid.

- Buses can be grouped into two main categories depending on clocking technique
- **Synchronous buses** have a "master clock"
 - Special line driven by a crystal oscillator (very precise timing)
 - Signal on the line is a square wave between 5 and 1GHz
 - All bus activities take a (fixed) integral number of these cycles (bus cycles)
- **Asynchronous buses** have no "master clock"
 - Bus cycles can be of any length
 - May differ between various pairs of devices
 - (bus still has a clock line)

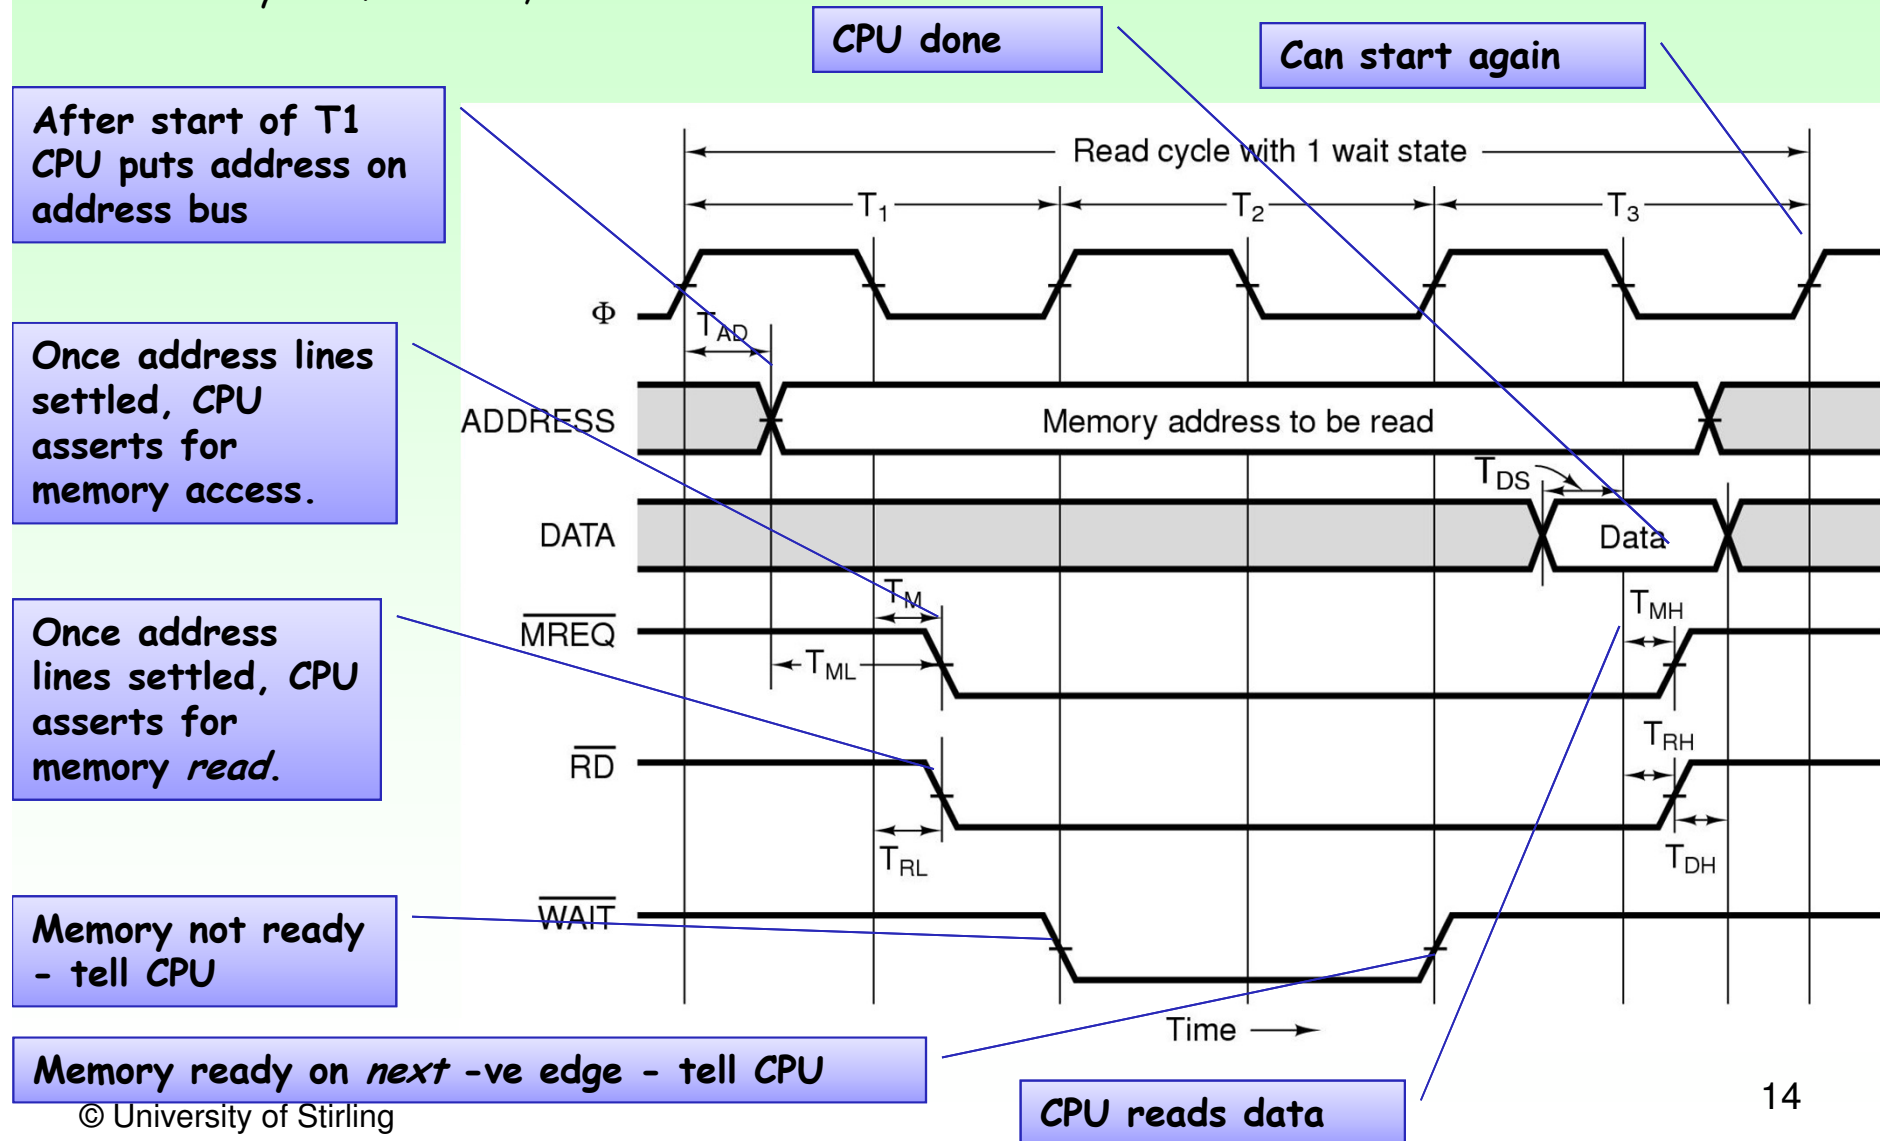
Synchronous Bus Transfer

- Figure shows one read and one write cycle.
- Each takes place in exactly one clock cycle.



A more complex example (Tanenbaum)

- 40MHz clock, - memory read: 40nsec
- bus cycle of 25 nsec, _____

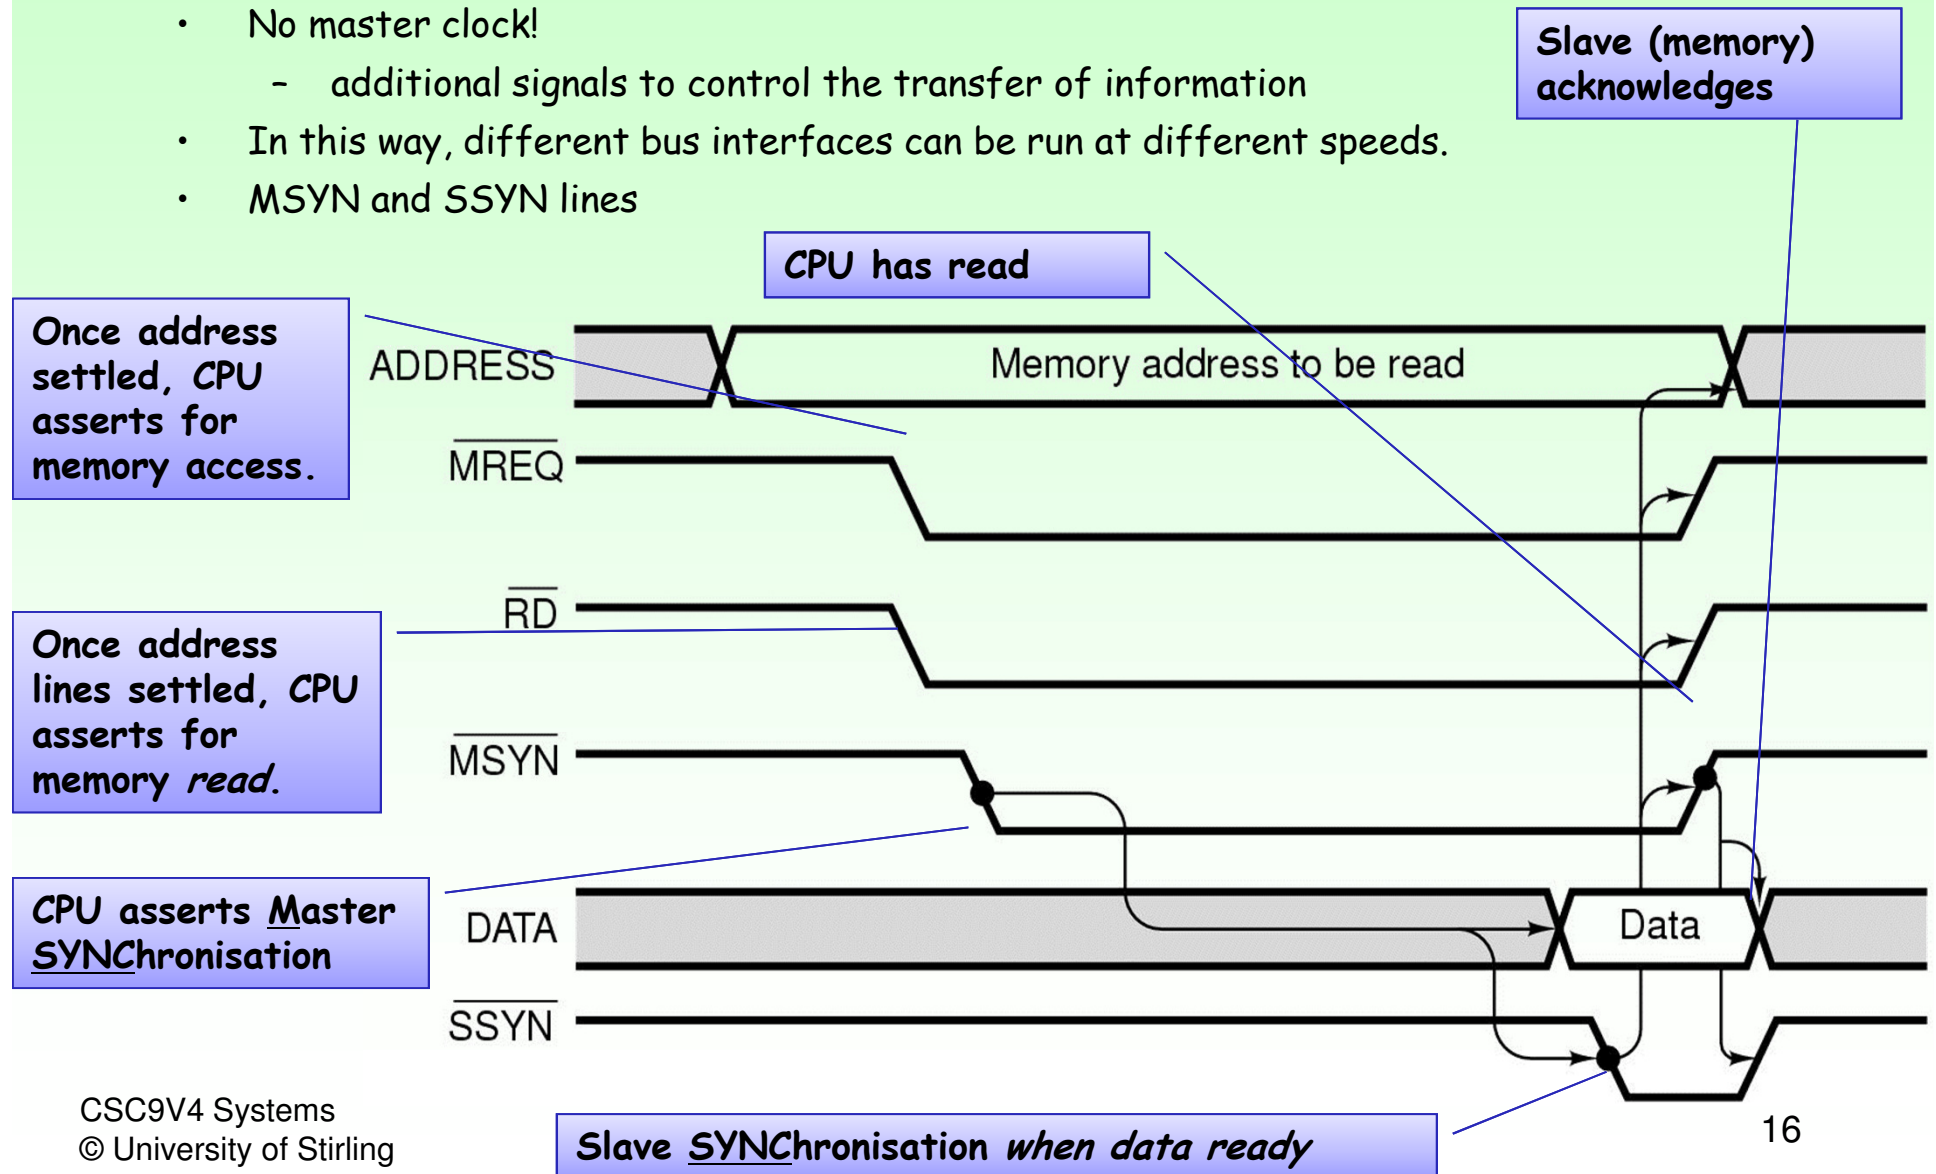


Problems of Synchronous Buses

- All devices on the bus must be capable of operating at the bus clock speed
- The bus needs to run at the speed of the slowest component on the bus
 - Implies wait states
- Making the whole system run rapidly then entails using fast components for everything on the bus.
- Everything works in multiples of the bus clock
 - If CPU and memory are capable of completing a transfer in 3.1 cycles, it is stretched to 4 cycles → no fractional cycles
 - Difficult to take advantage of future improvements of technology
 - What if new memory with an access time of 20 nsec becomes available?
 - Or with 10 nsec?
 - No good handling of heterogeneous collection of devices
 - Good PC hardware designers account for this with memory and device selection to maximise throughput

Asynchronous Buses

- No master clock!
 - additional signals to control the transfer of information
- In this way, different bus interfaces can be run at different speeds.
- MSYN and SSYN lines



Full Handshake

- A set of signals that interlocks as MSYN and SSYN
→ Full Handshake
- MSYN is asserted
- SSYN is asserted in response to MSYN
- MSYN is negated in response to SSYN
- SSYN is negated in response to negation of MSYN
- Full handshakes are timing independent
 - Each event is caused by a prior event
 - No clock pulses involved
 - If a master-slave pair is slow, other master-slave pairs are not affected

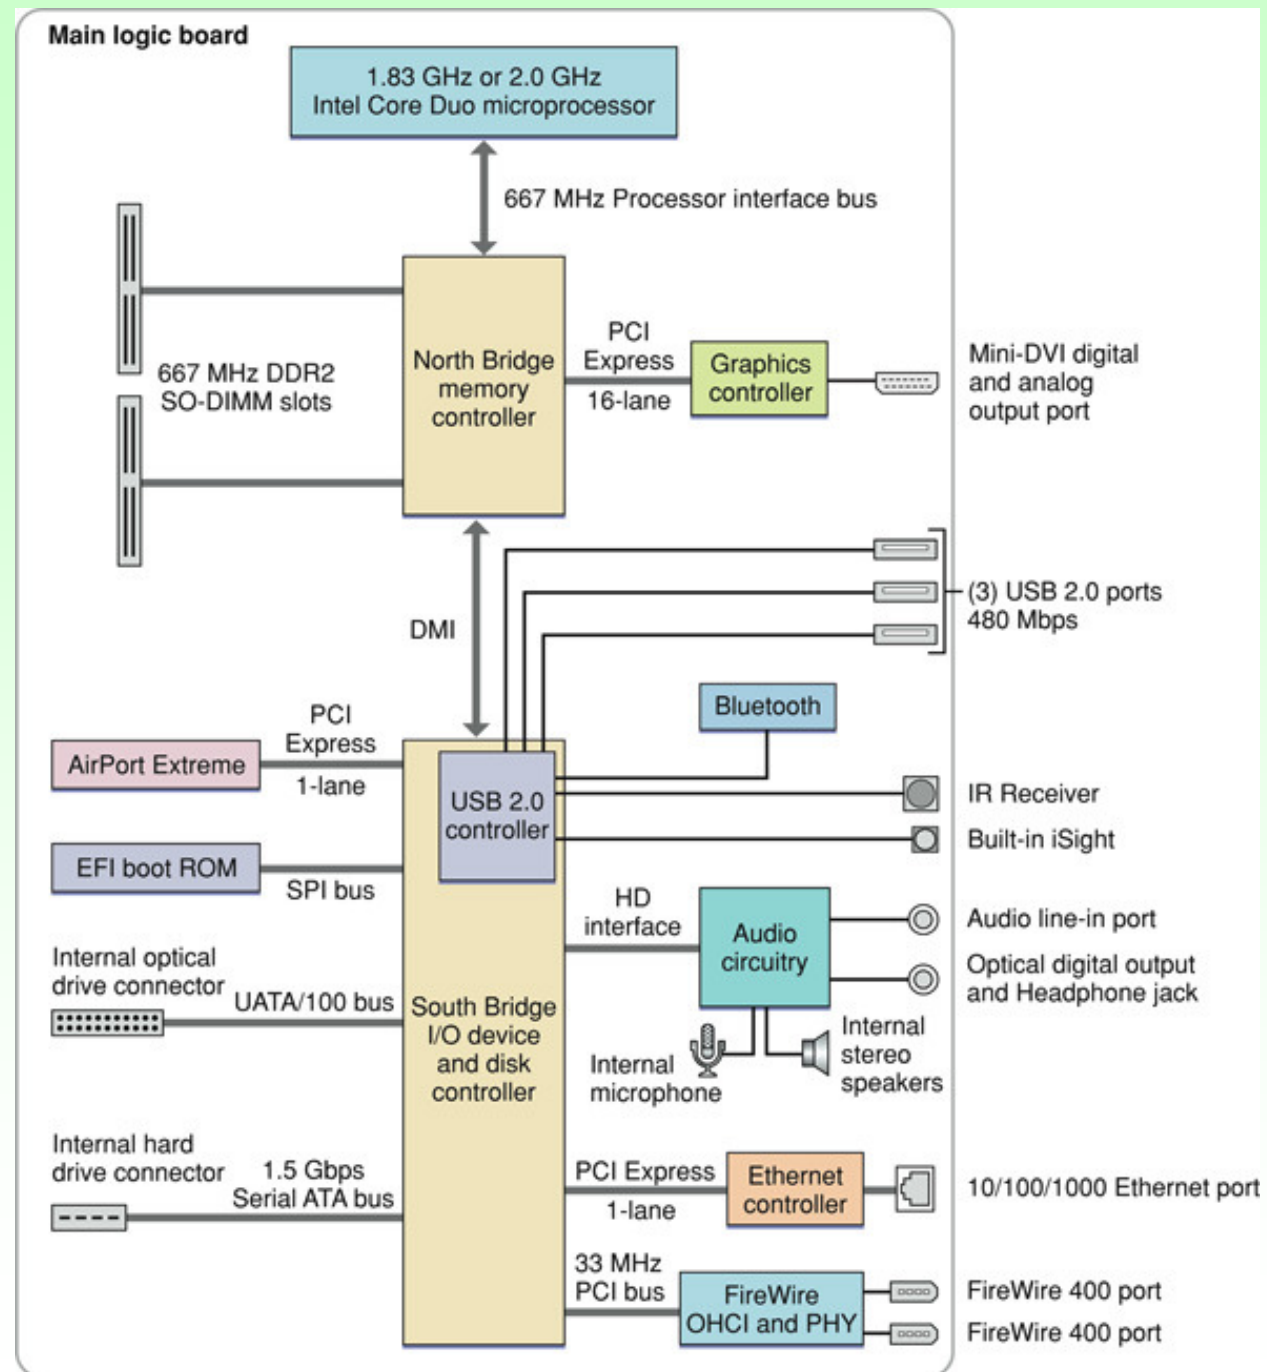
Why Asynchronous Buses?

- Asynchronous bus transfer allows devices which run at different speeds to share one bus without a speed penalty.
- The slave can spend a variable amount of time
 - accessing/retrieving data
- Similarly, the master can spend a variable length of time
- Actions need not occur at every clock pulse.
- However, most buses are synchronous buses
 - Easier to build; effective if components are chosen properly
 - Can use multiple busses to cope with fast and slow subsystems
 - Lot of investment in synchronous bus technology

Buses in Modern PCs

- Requirements
 - very fast transfer between CPU and Memory
 - fast transfer between fast peripherals and CPU/Main Memory
 - slower transfer for slower peripherals
 - slow and inexpensive transfers for slow peripherals
- CPU/Memory
 - usually a proprietary bus: very high speed
- PCI bus: High speed bus, wider and faster than ISA
- PCIe bus: high speed serial link
- USB bus: Generic 'one size fits all' (mice, keyboards, printers)
- ISA bus (out of date): Compatibility reasons

- Modern bus architecture



The PCI Bus


- Need for high speed - consider displaying a video: huge data rates required
- PCI = Peripheral Component Interconnect
- designed by Intel in 1990, open access: public domain (used by Apple as well)
- PCI transfers 32-bit or 64 bit data at 33 MHz or 66 MHz or 132 MHz
- Provides a high bandwidth link to fast peripherals.
- Bus is synchronous (mostly...)
- Data and address lines are multiplexed
- Bus has facilities for transferring bus mastership in each cycle
 - usually requires 1 cycle between reads and writes
- Bus has also facilities for burst mode transfers
- Bus also has interrupt lines

Other buses

- PCI Express: replacement for PCI, but serial, faster and full duplex
- AGP: bus used purely for graphics devices (usually just 1 graphics card)
- Buses for discs
 - SCSI
 - EIDE
 - ATA, SATA
- Serial buses
 - FireWire
 - USB
 - Thunderbolt

PCI Express

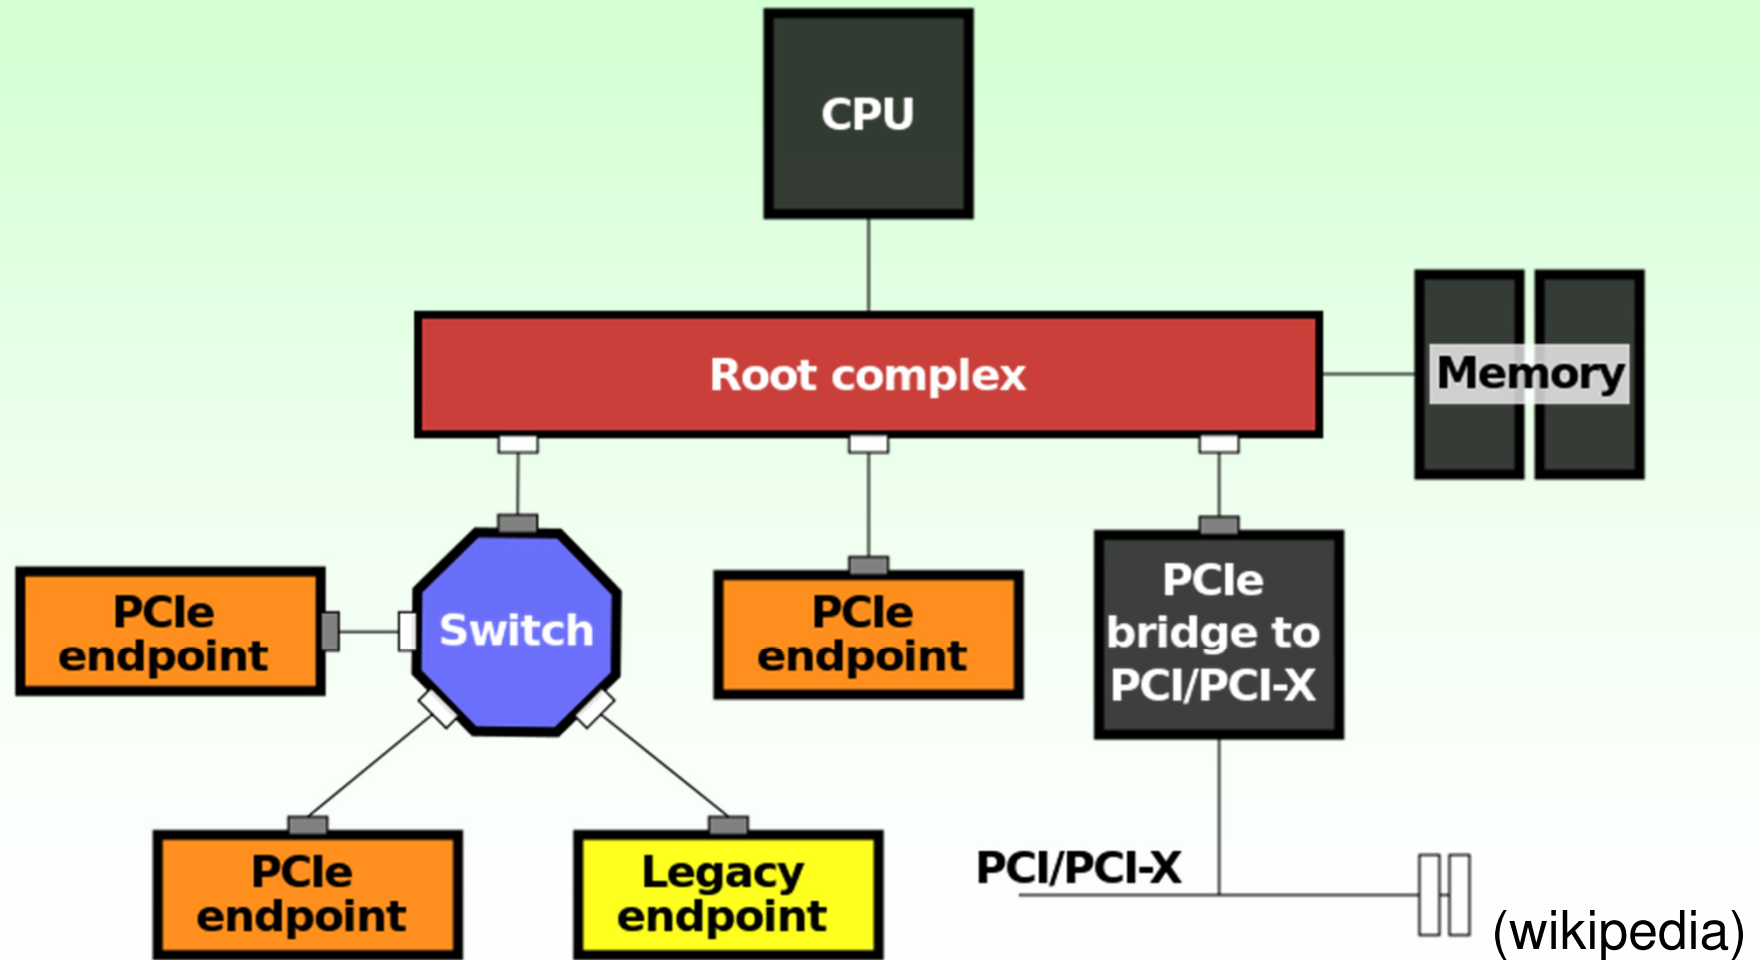
- PCIe is a serial bus that has a smaller footprint and is faster than PCI



(wikipedia)

PCI Express

- Point-to-point connections with packet-switching (like internet)



End of Lecture