

Graphical User Interfaces	Higher-level Programming
Operating Systems	
Low-level Programming	
Basic Machine Architecture	
Silicon	

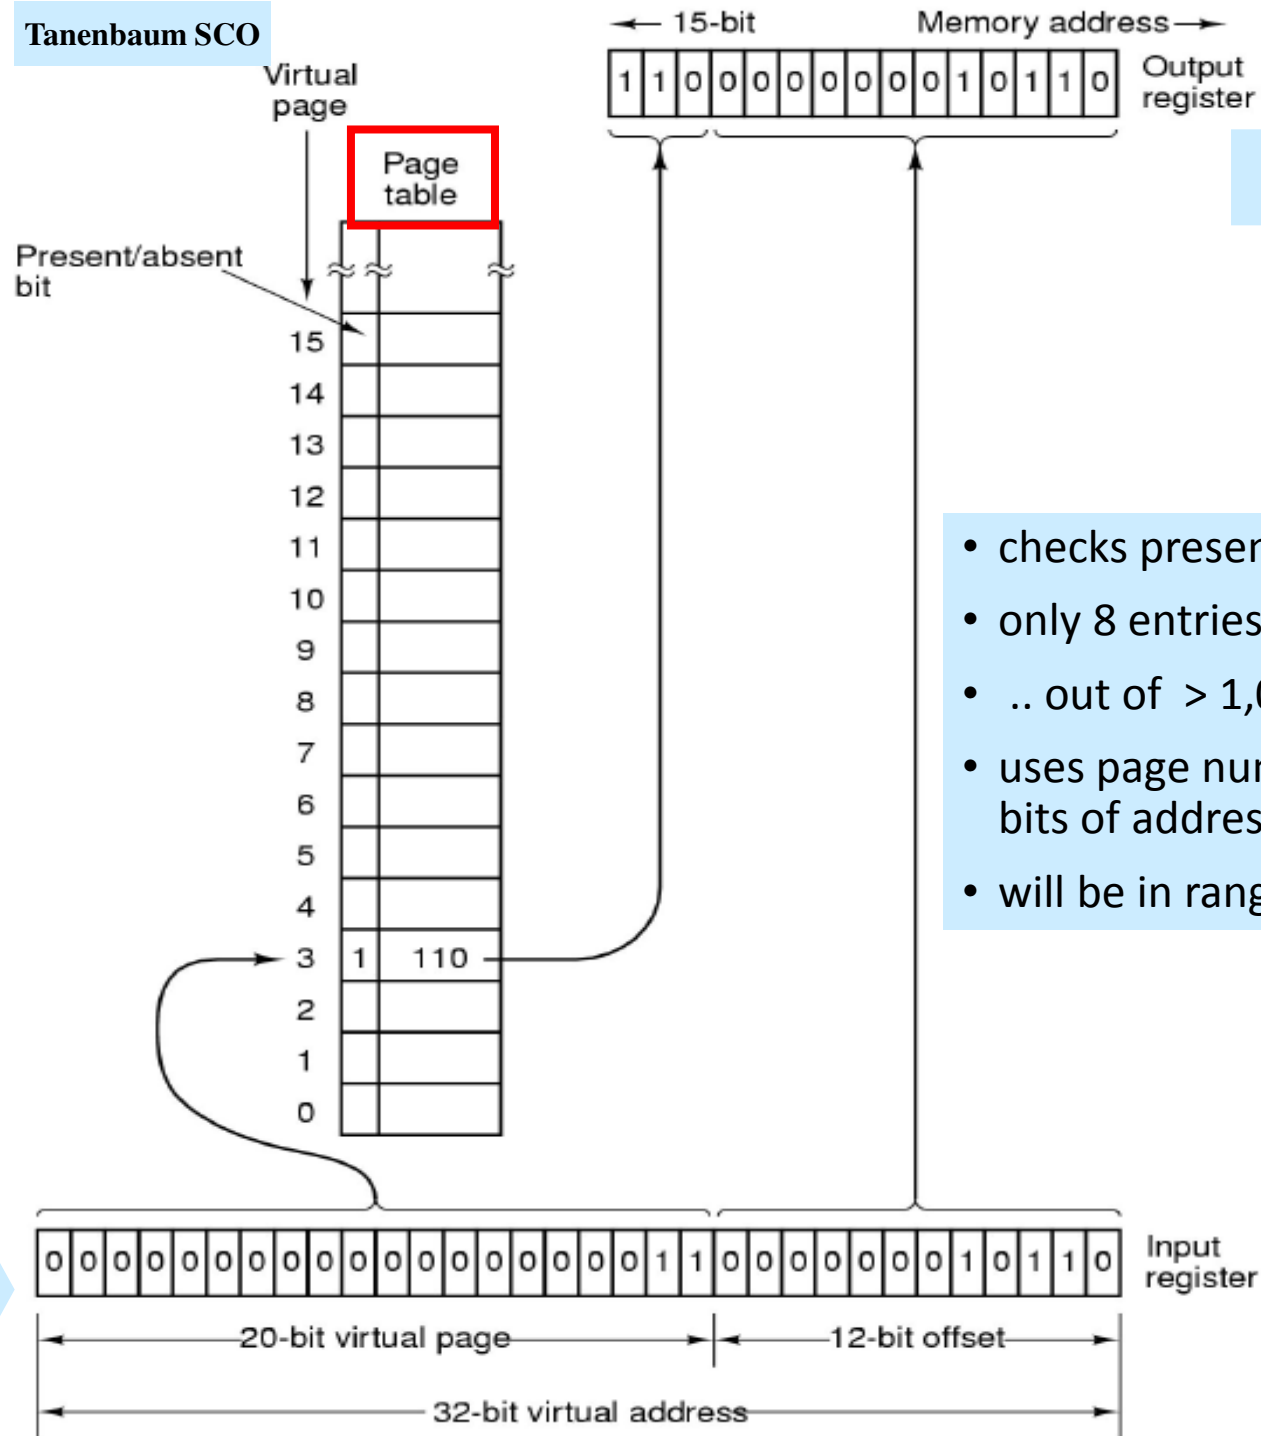
# CSCU9V4 Systems

Systems lecture 15  
Operating Systems 2

**VM issues: Page sizes, Segmentation**

implemented on a MMU

# Tanenbaum SCO



15 bit P.A.

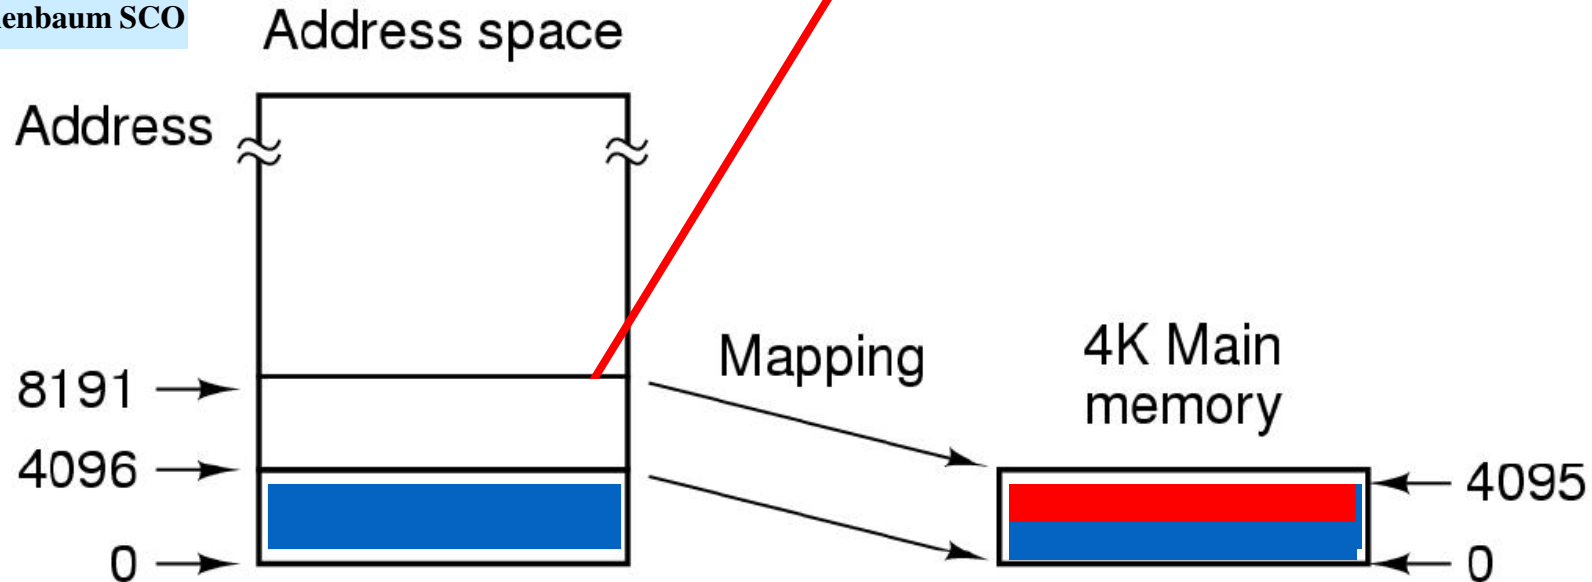
- checks presence bit
- only 8 entries set
- .. out of > 1,000,000
- uses page number for top bits of address
- will be in range 0 to 7

# What size should a page be?

- Fragmentation:
- assume a 4K pages
  - ... so that is the finest resolution we can use
- ... if we have a 5K program
  - ... must 'use' 8K
- **internal fragmentation**

• 3 K 'wasted'

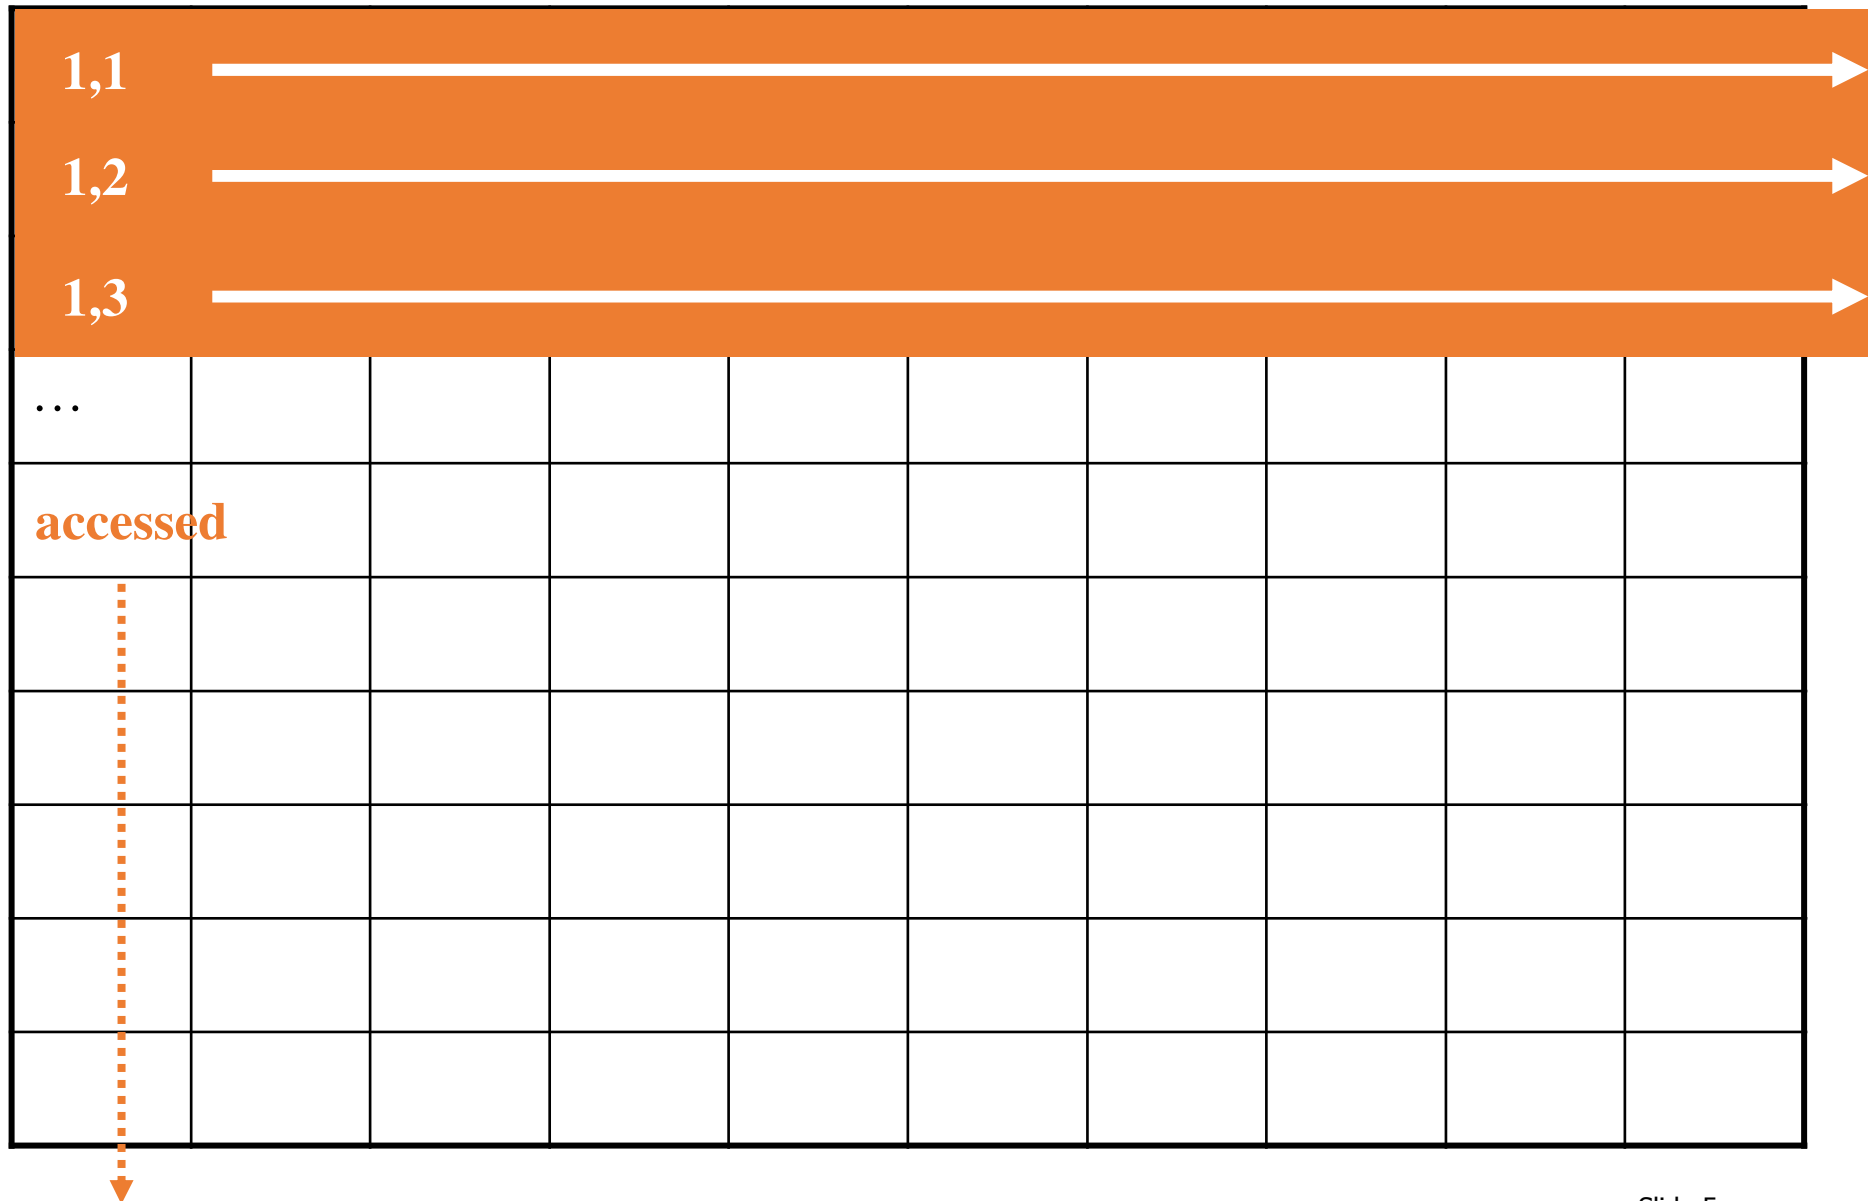
Tanenbaum SCO



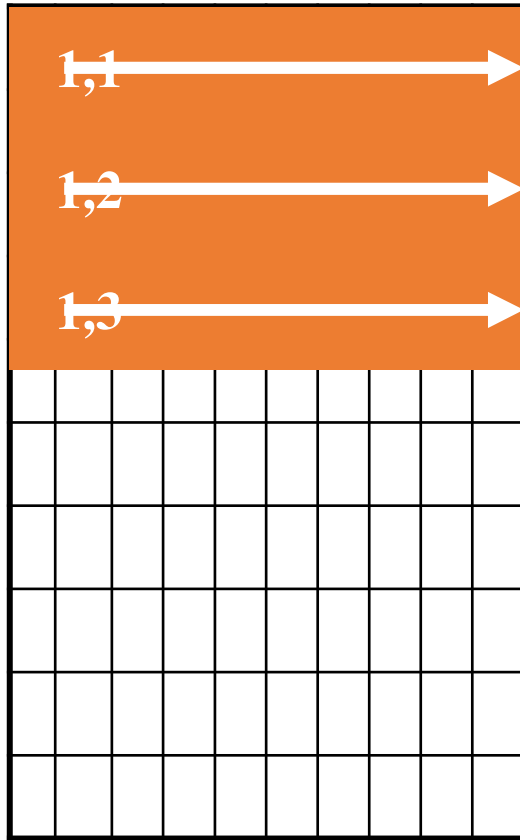
# What size should a page be?

- smaller page size wastes less space.
- if the page size is  $N$  bytes
  - the average wastage  $N/2$  bytes
  - (Of course, we don't actually waste nearly this much: mostly the blocks of virtual memory we bring into physical memory are much more than 1 page in size: we only waste this much of the last page)
- [probability of wasted space is the same for  $1..(N-1)$  ]
- ... so best to have page sizes small ?
- Also...
- consider a 10,000 by 10,000 matrix of **8-byte floats**  $A[10000, 10000]$
- say bytes of  $A$  stored in **column** order  $A[1,1], A[2,1], A[3,1] \dots$

What size should a page be (3)



## What size should a page be? (4)



- Let's say that we carry out row-ordered calculations, (i.e.  $A[1,1]$ ,  $A[1,2]$ ,  $A[1,3]$  ....) inside a loop executed a few hundred times

- the elements are 80,000 bytes apart
- (i.e. row elements separated by 79,992 bytes!)



## What size should a page be? (5)

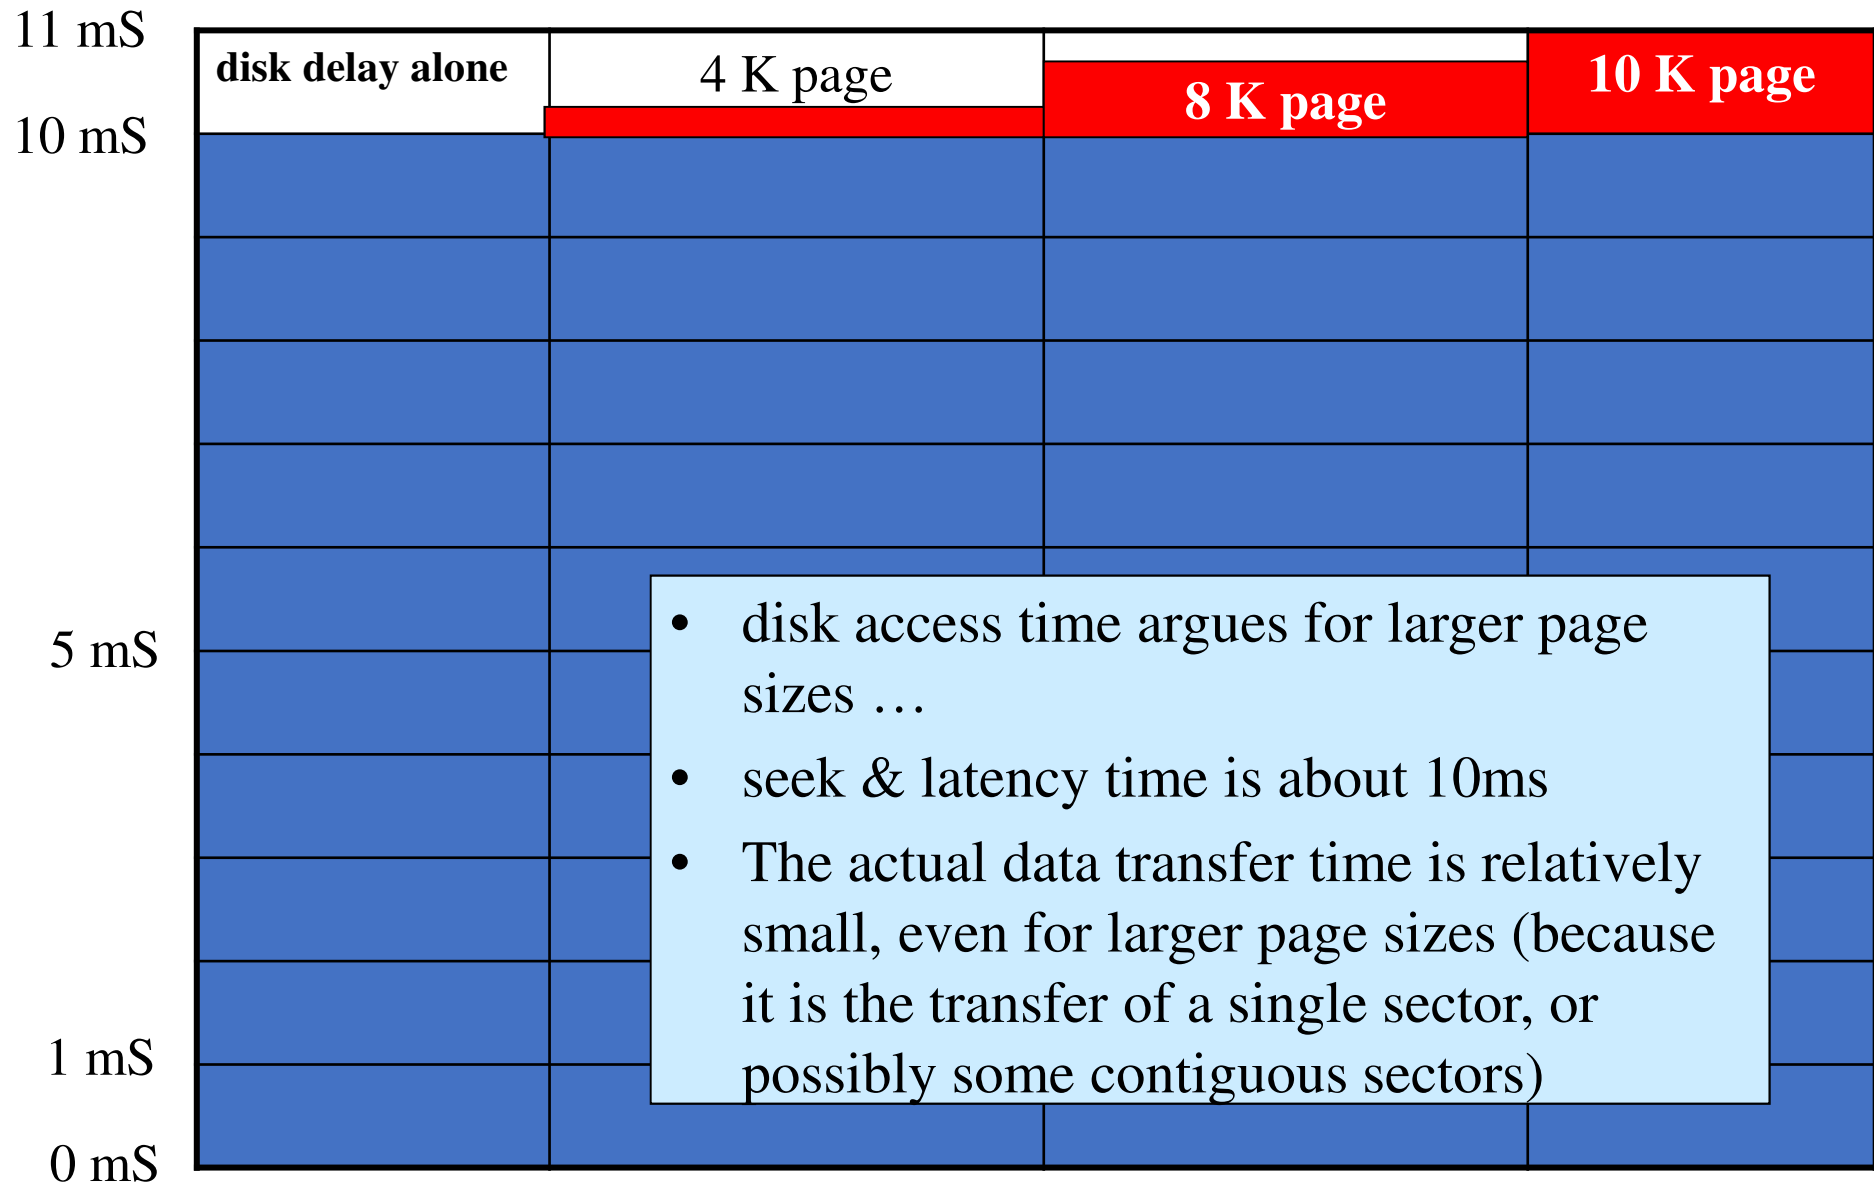
- each element down a column requires a page to be brought in to physical memory from virtual memory
  - so with 8K pages we require  $10,000 \times 8K = 80\text{Mbytes}$  (approx)
  - i.e. the working set would be 80Mbytes
- if the page size was 1K, only 10Mbytes would be needed.
- so a 64M byte physical memory will **thrash** with 8K pages
- ... but not with 1K pages!
  - ... so again smaller pages seem better
- (How realistic is this? Huge array – 10,000,000,000 elements long)
  - .... **but** ..

## What size should a page be (6)

- There's a question of the efficiency of moving pages from virtual to physical memory.
- disk access time argues for larger page sizes
- seek & latency time is about 10ms
- data rates from disk controller are above 10Mbytes per second
- ... or better than 1mS for each 10Kbytes
- Further, efficiency for disc reads/writes suggests that sectors on disk be quite large
  - Recall whole sectors are read or written at a time
- And efficient transfer between disc/main memory suggests page size be the same as sector size.



## What size should a page be (7)

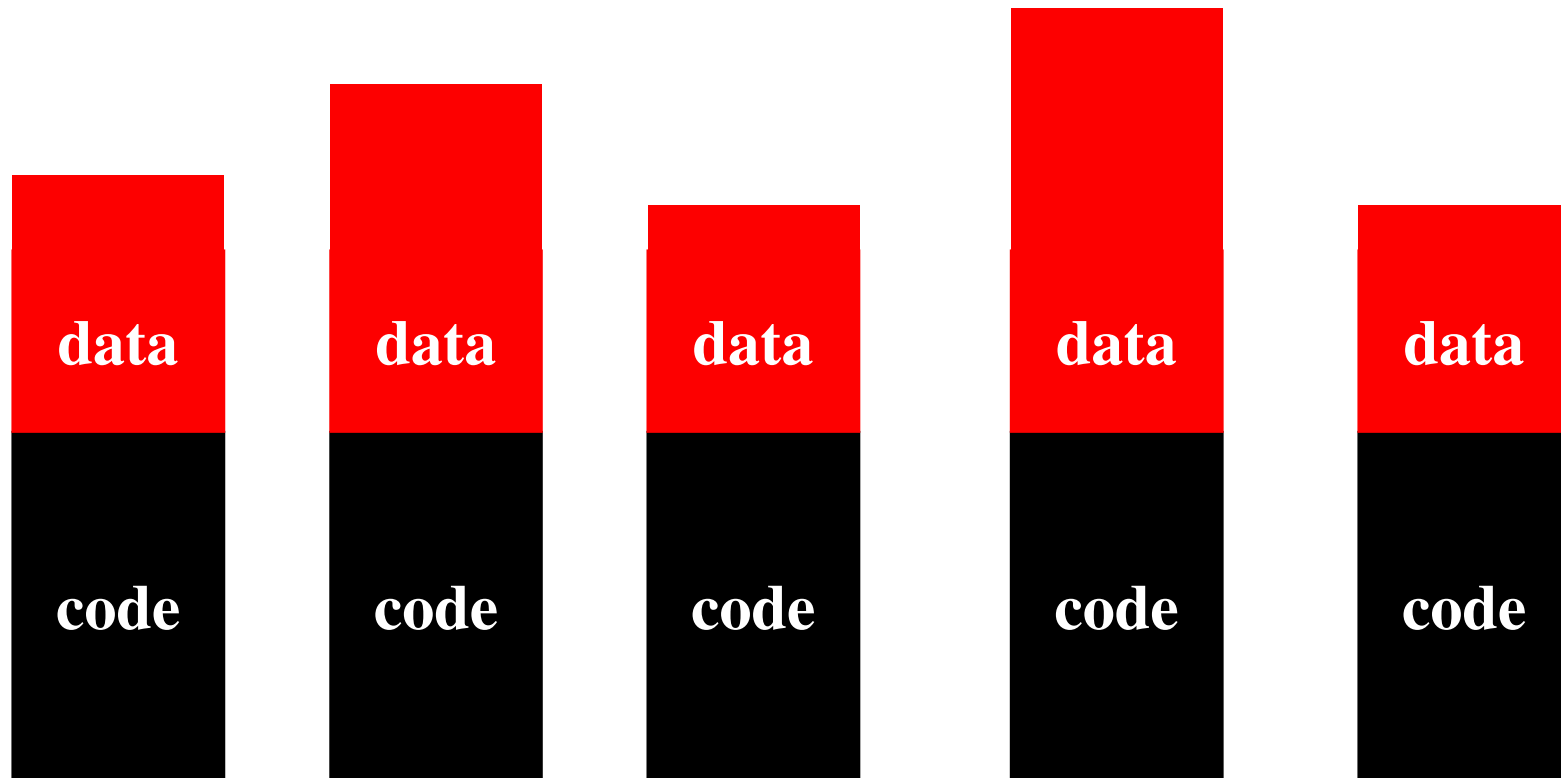


# Virtual Memory: Review.

- Virtual memory is now virtually universal for desk top machines
- Users do not need to worry about the total size of their programs
  - They can just keep starting up more, (so long as they fit in virtual memory).
- Efficient transfer between disk & memory argues for large page sizes
- Efficient memory use argues for smaller page sizes
  - IA64 architecture allowed for 4K or 4M pages!
- For a program to run efficiently, it is working set size that matters
- When the working set size for the programs in use exceeds the physical memory, the system runs very slowly indeed.
  - Thrashing occurs

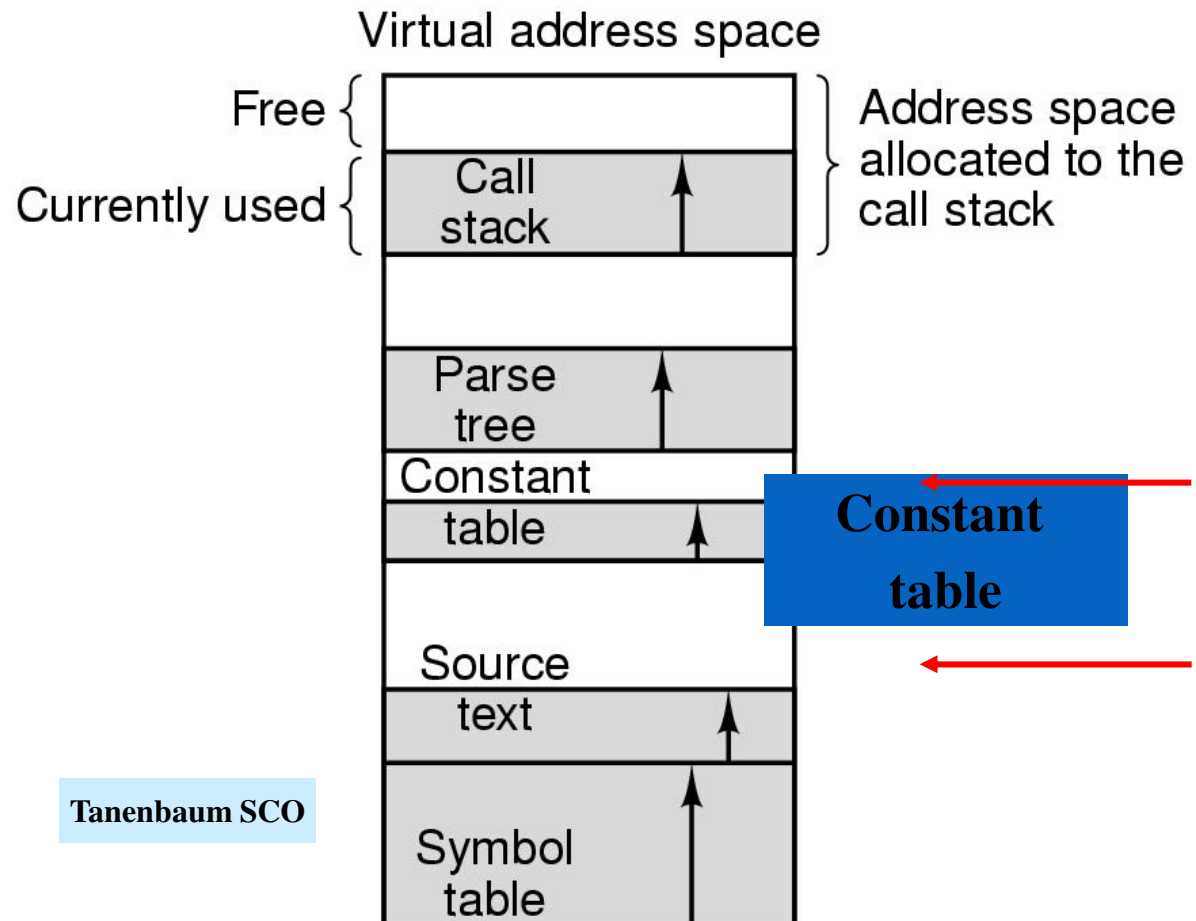
## Segmentation (1)

- what if the a program has a dynamic volume of data as it runs?
- Program has code and data



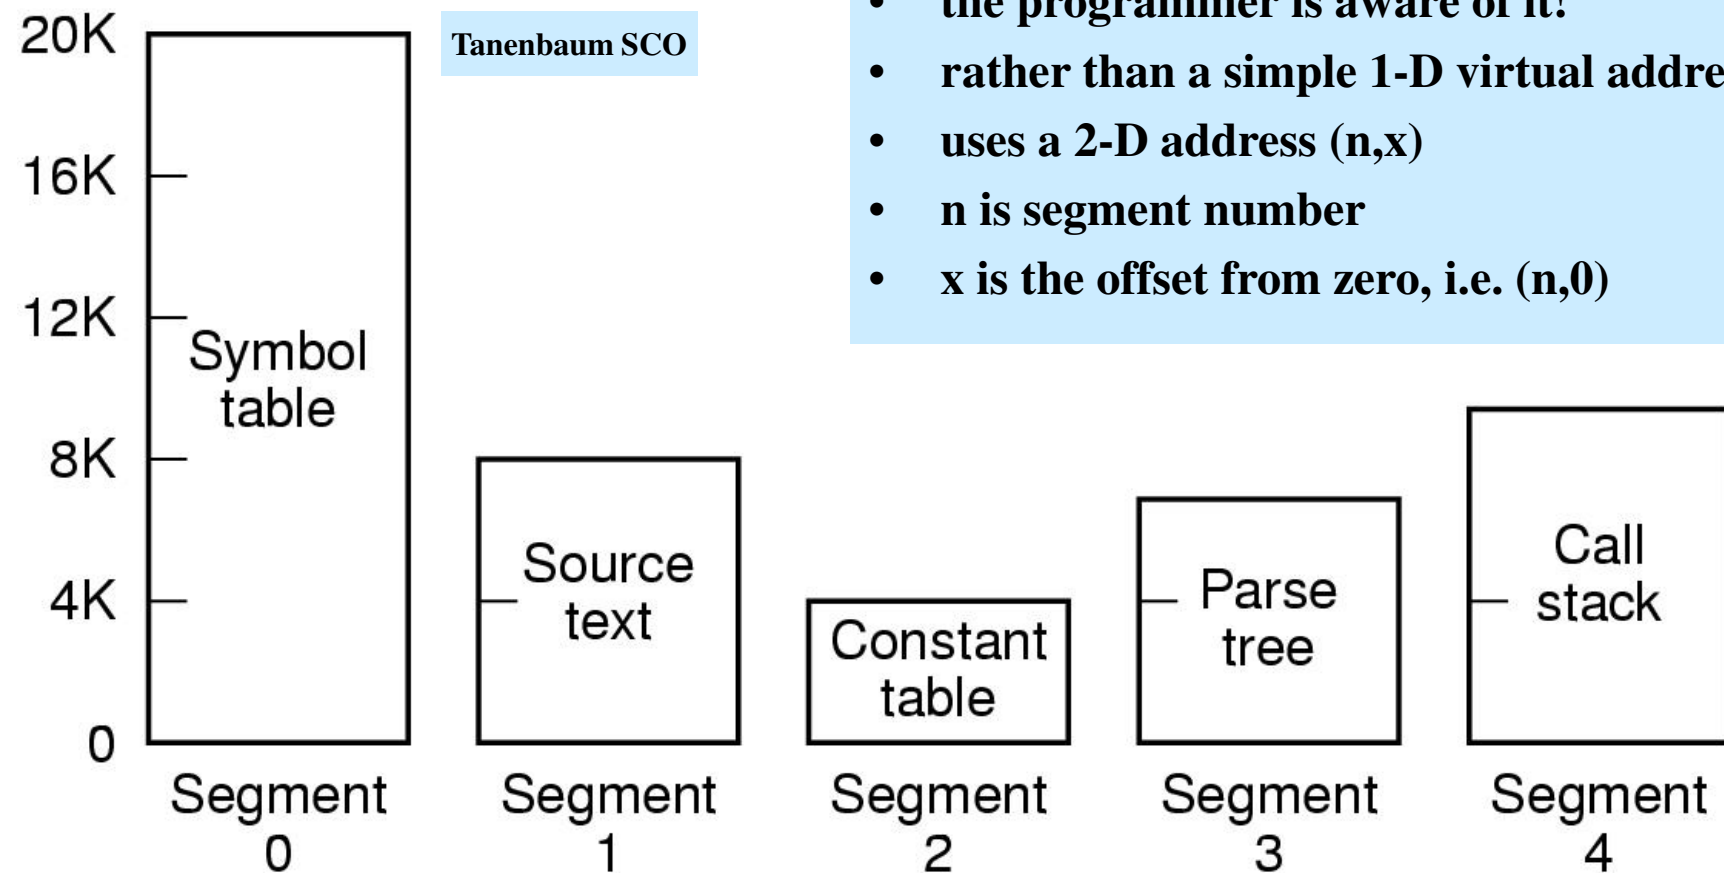
## Segmentation (2)

- if the data that is changing across a number of tables ....
- for example a compiler
- Or a program creating and destroying objects
- Or a program calling methods (or functions) which have a lot of local data



## Segmentation (3)

- better to split the data up ....
- ... into segments ...



- **each segment starts at zero**
- **to some max value**
- **shrink & grow independently**
- **the programmer is aware of it!**
- **rather than a simple 1-D virtual address**
- **uses a 2-D address (n,x)**
- **n is segment number**
- **x is the offset from zero, i.e. (n,0)**

## Segmentation (4): *more advantages (1)*

- other advantages ...
  - procedure linking is improved
  - can be compiled **separately** – all starting at address zero
  - do not need to compile the complete program

**procedure A**

**procedure B**

**procedure C**

**procedure D**

## Segmentation (5): *more advantages (2)*

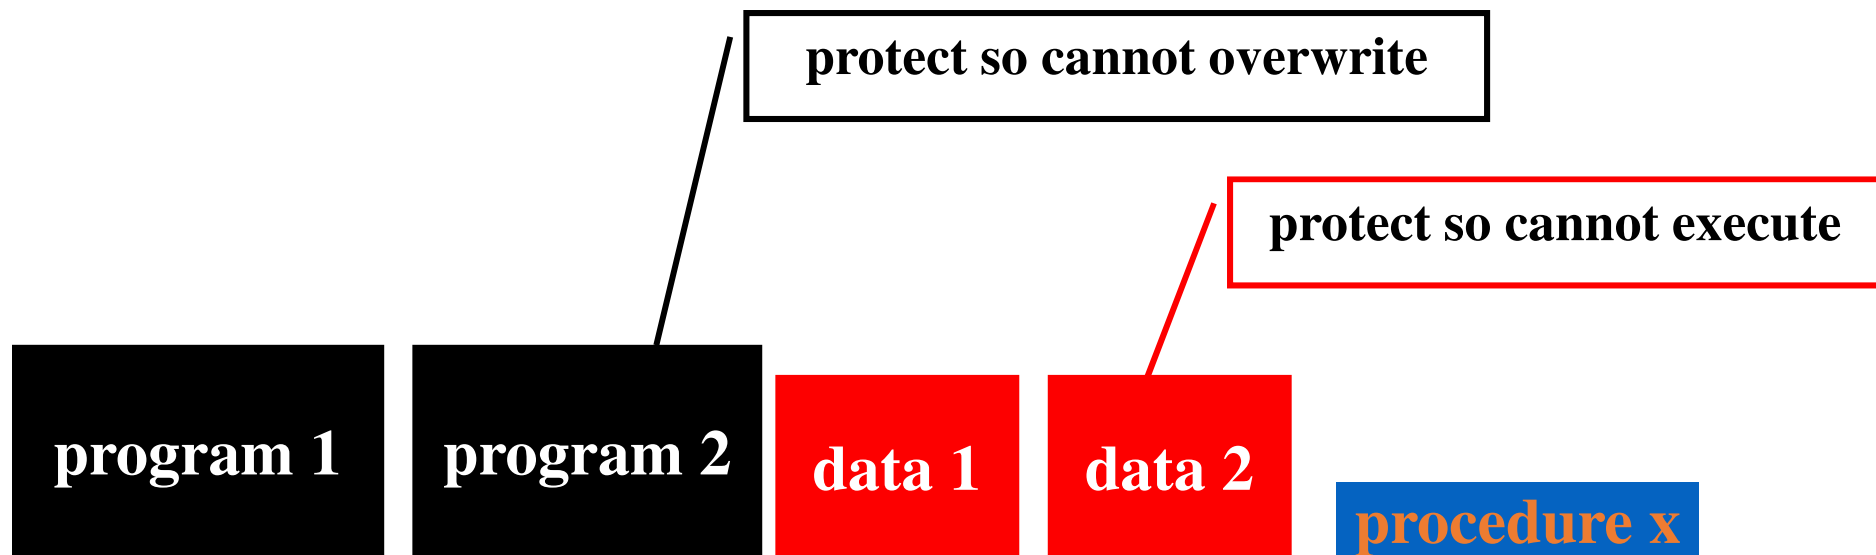
- other advantages ...
  - can also share procedures between processes



## Segmentation (6): *protection*

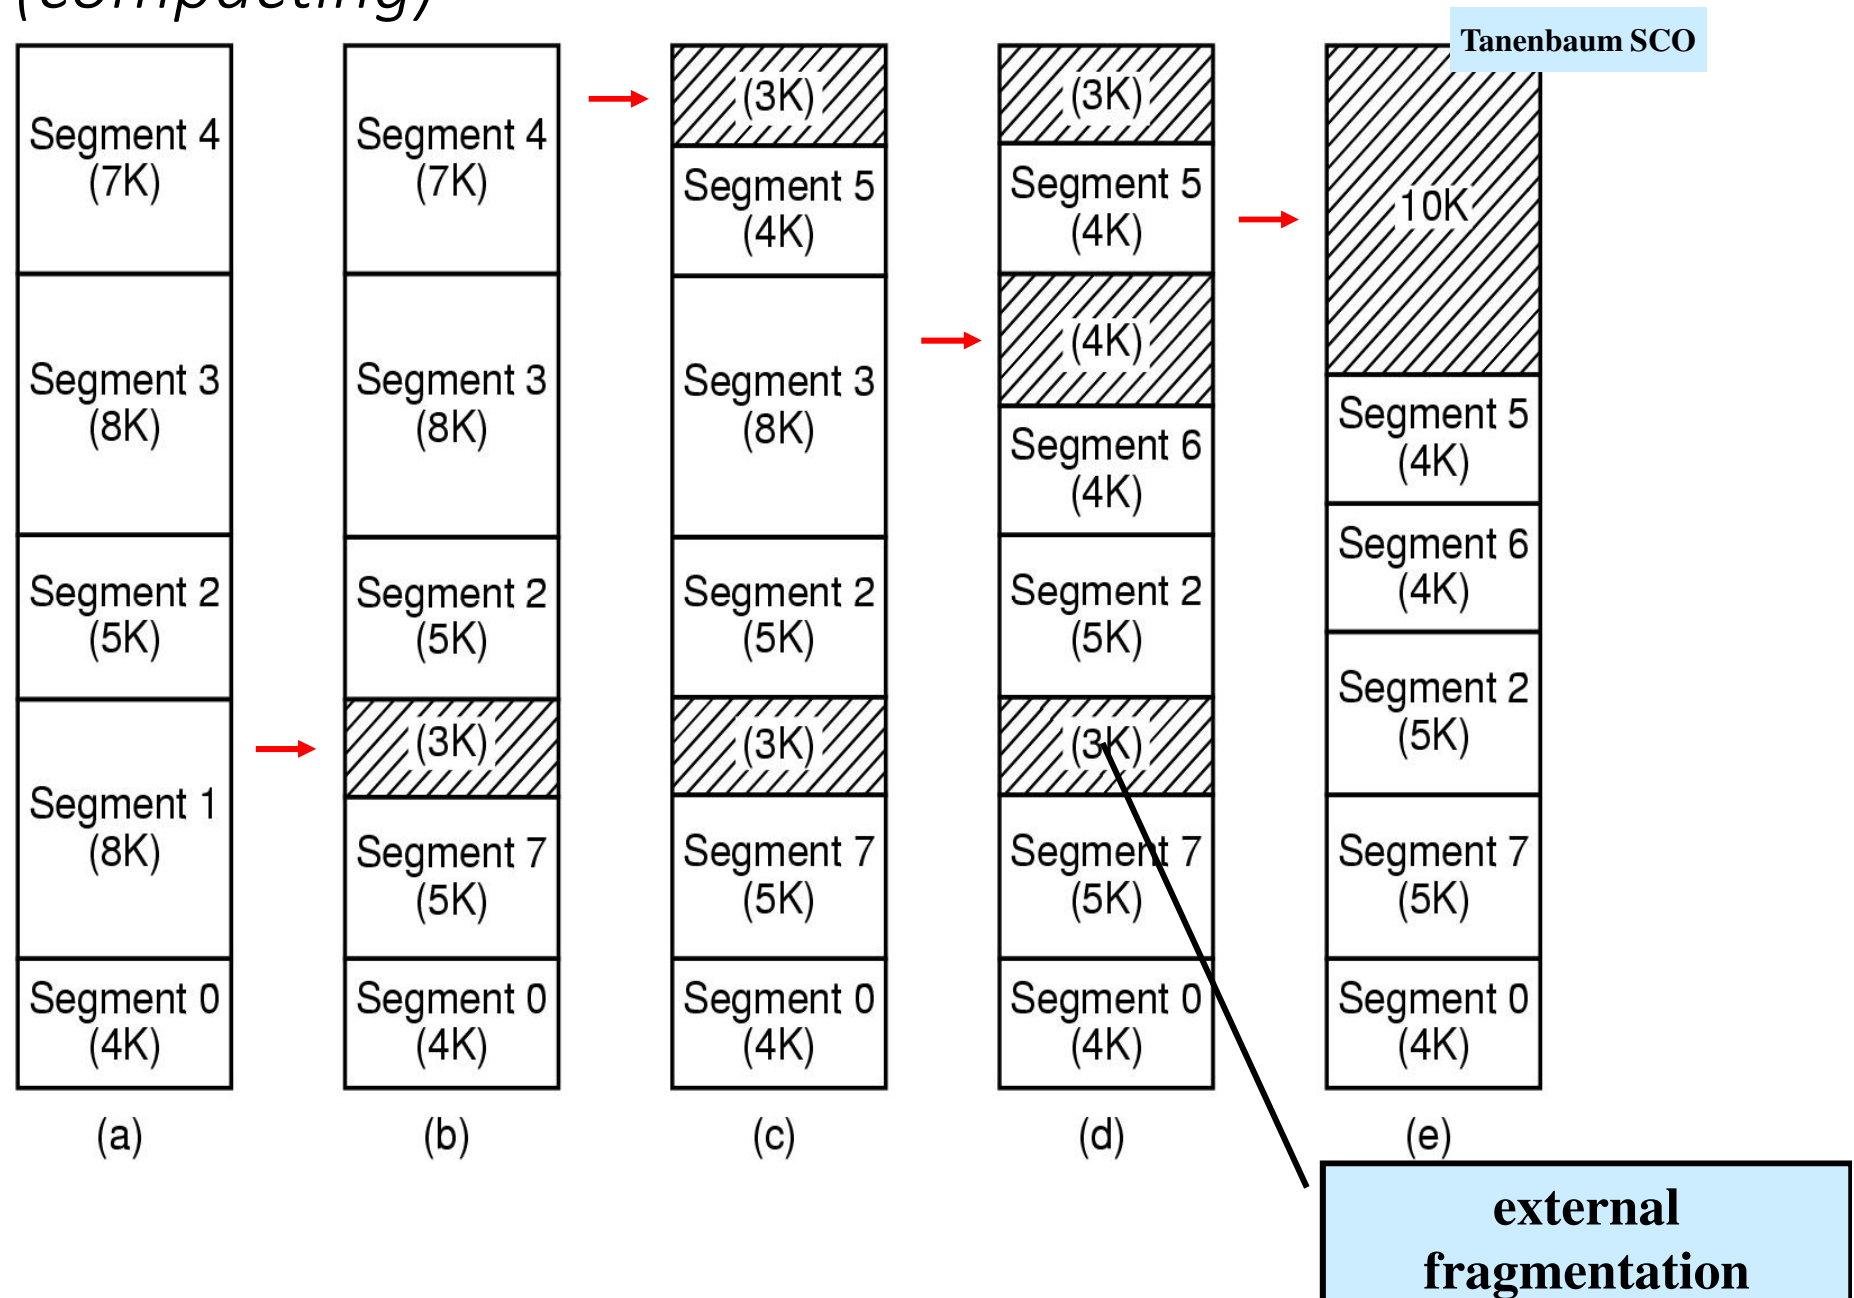
- yet another advantage ...

- the programmer is aware of each segment ... so ...
- they can be of one type
- so it can be *protected* as say read/write, and any attempt to execute it will fail ... useful if program jumps to an incorrect address
- And program can be protected as execute only, so any attempt to overwrite it will cause a fault.

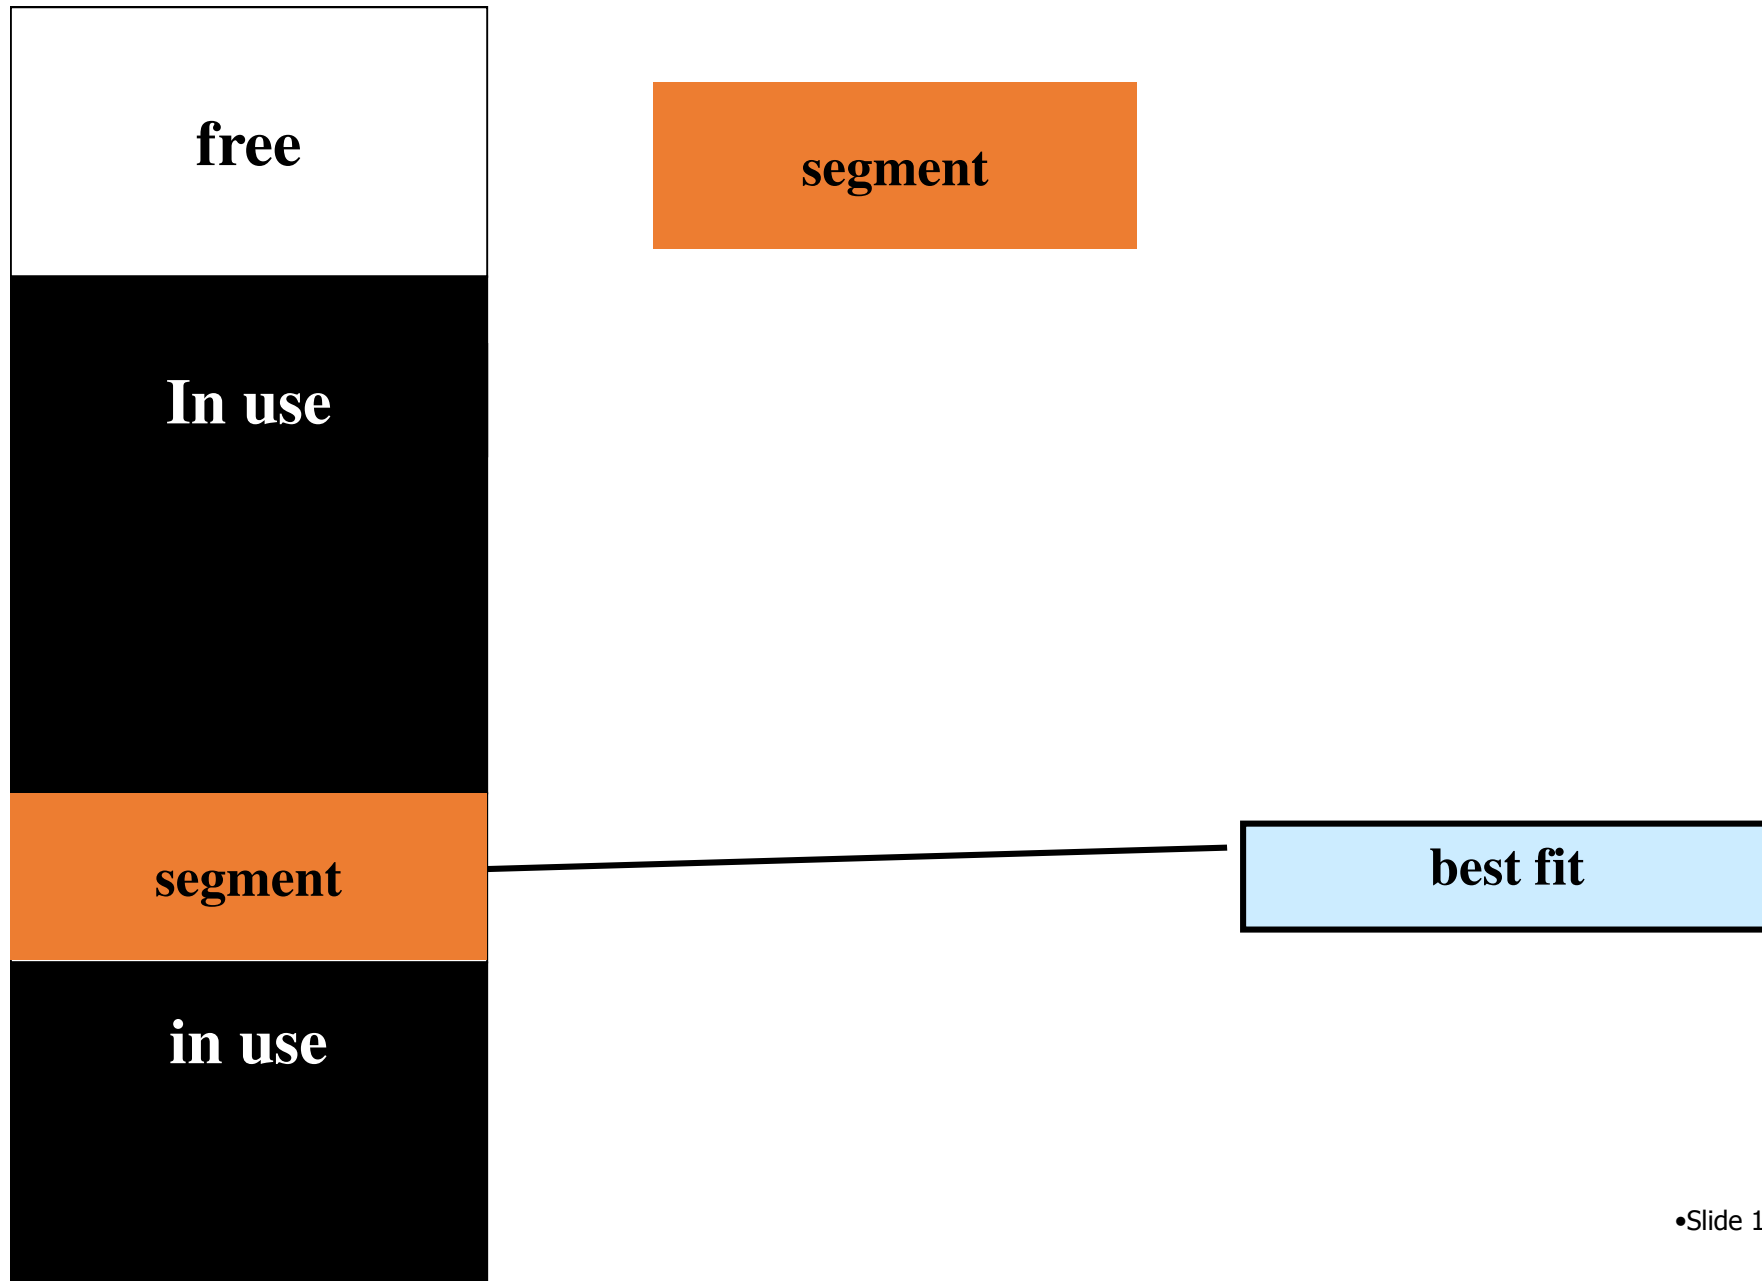




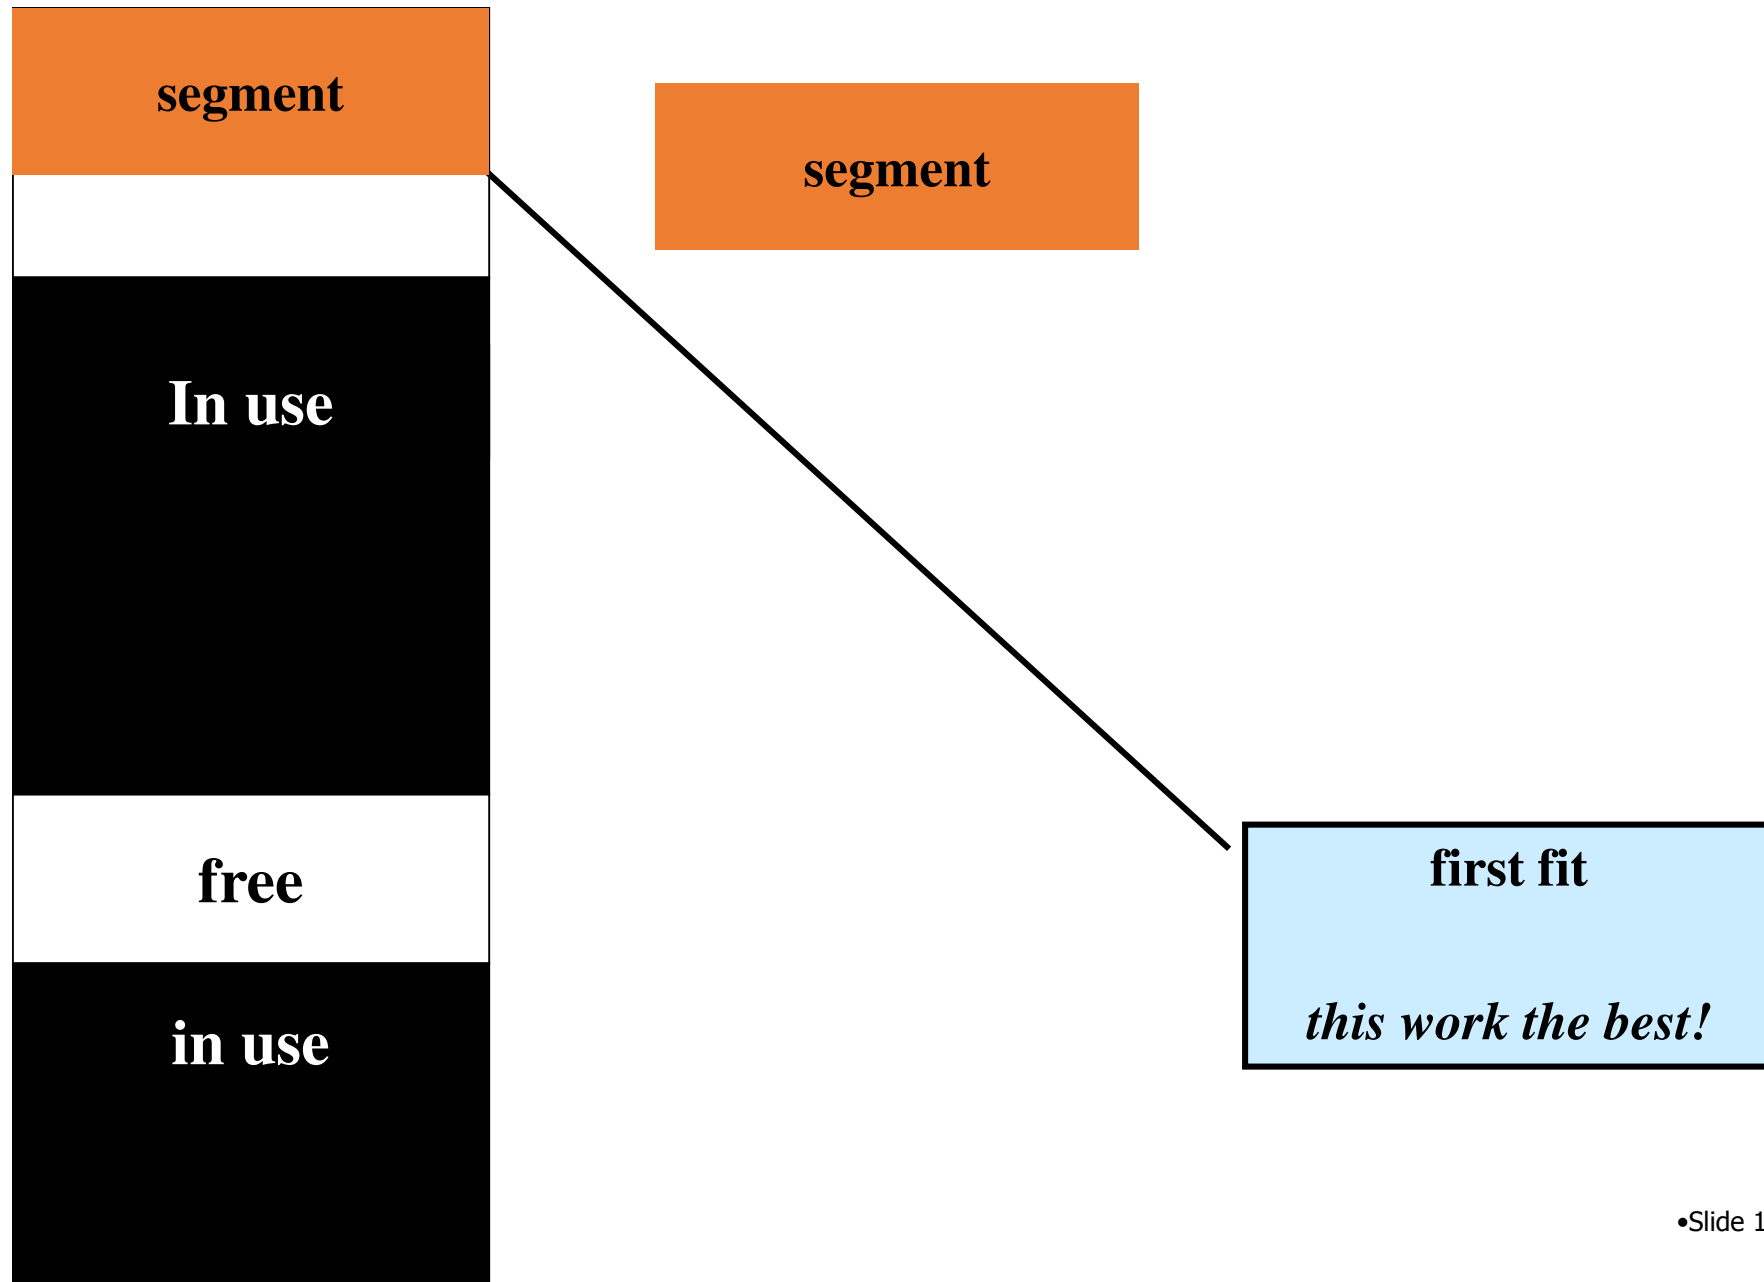
# Segmentation (7): non-paged *implementation* (*compacting*)



## Segmentation (8): *implementation (if no time to compact)*

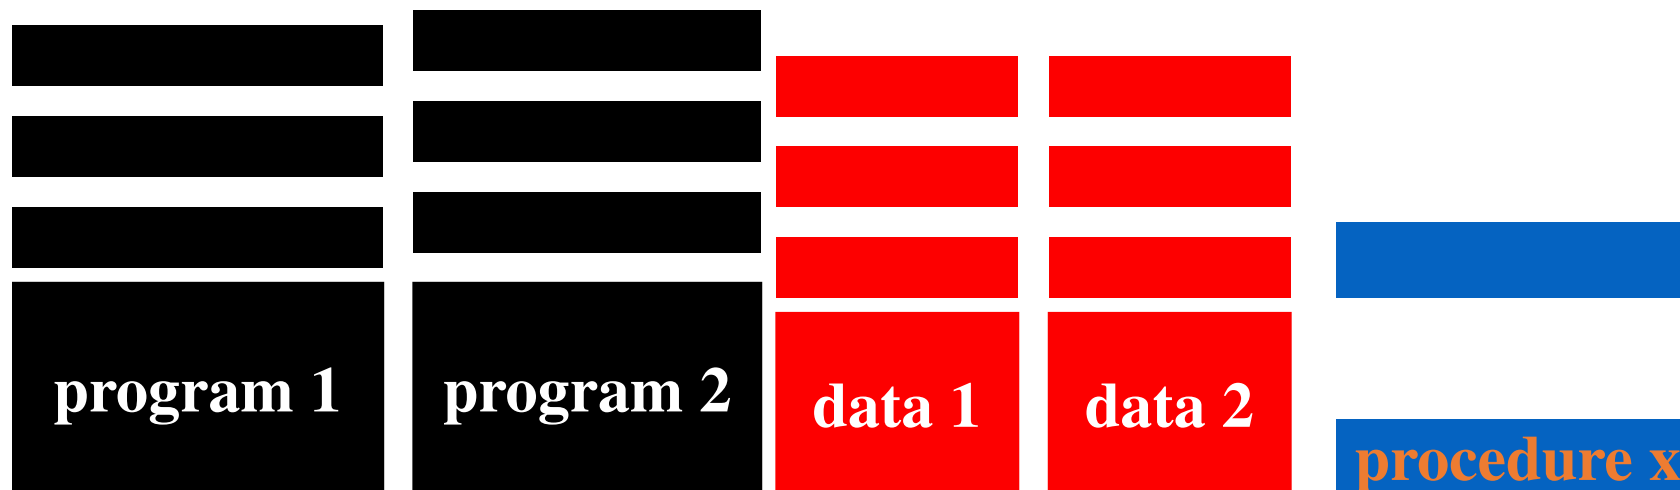


## Segmentation (8): *implementation (if no time to compact)*



## Segmentation (9): *implementation (paging)*

- Alternatively - each segment can be split into pages
  - each segment is a 1-D memory space
  - exactly as we saw before
  - so we can apply the same techniques as before
- common to have segmented paging arrangements
  - Intel
  - still a 2-D view to programmer – paging is transparent (as before)



End of lecture